

# Bayesian Optimization Hackathon for Chemistry and Materials

Sterling G. Baird,<sup>1,2</sup> Mehrad Ansari,<sup>2</sup> Gage Erickson,<sup>3</sup> Zartashia Afzal,<sup>4</sup> Qianxiang Ai,<sup>5</sup> Alexander Al-Feghali,<sup>6</sup> Mathieu Alain,<sup>7</sup> Matias Altamirano,<sup>8</sup> Thomas Andrews,<sup>9</sup> Andy S. Anker,<sup>10,11</sup> Rija Ansari,<sup>12</sup> Samuel Ampofo Appiah,<sup>13</sup> Raul Astudillo,<sup>14</sup> Ruhana Azam,<sup>15</sup> Mohammed Azzouzi,<sup>16</sup> Suneel Kumar BVS,<sup>17</sup> Ben Blaiszik,<sup>18</sup> Anna Borisova,<sup>19</sup> Andres Bran,<sup>19</sup> Pengfei Cai,<sup>5</sup> Ting-Yeh Chen,<sup>20</sup> Curtis Chong,<sup>21</sup> Samantha Corapi,<sup>21</sup> Mark Croxall,<sup>21</sup> Gbetondji Dovonon,<sup>7</sup> Jose Manuel Napoles Duarte,<sup>22</sup> Andrew Falkowski,<sup>22</sup> Giuseppe Fisicaro,<sup>23</sup> Martin Fitzner,<sup>24</sup> Quinn Gallagher,<sup>25</sup> Sabah Gaznaghi,<sup>26</sup> Jérôme Genzling,<sup>6</sup> Christoph Griebel,<sup>27</sup> Ryan-Rhys Griffiths,<sup>28</sup> Taicheng Guo,<sup>29</sup> Kehan Guo,<sup>29</sup> Nipun Gupta,<sup>2</sup> Ankur Gupta,<sup>30</sup> Mohammad Haddadnia,<sup>31</sup> Yuyang Han,<sup>32</sup> Joscha Hoche,<sup>24</sup> Alexander V. Hopp,<sup>24</sup> Marko Huang,<sup>21</sup> Ayodeji Ijishakin,<sup>33</sup> Ramsey Issa,<sup>22</sup> Yeonghun Kang,<sup>34</sup> Jungtaek Kim,<sup>35</sup> Akshay Kudva,<sup>20</sup> Rubén Laplaza,<sup>19</sup> Magdalena Lederbauer,<sup>19</sup> Shi Xuan Leong,<sup>21</sup> Paul W. Leu,<sup>35</sup> Viola Muning Li,<sup>24</sup> Mingxuan Li,<sup>35</sup> Tao Liu,<sup>6</sup> Stanley Lo,<sup>21</sup> Jakub Lála,<sup>16</sup> Osman Mamun,<sup>36</sup> Owen Melville,<sup>2</sup> Michail Mitsakis,<sup>37</sup> Cameron Movassaghi,<sup>32</sup> Madhav Reddy Muthyala,<sup>20</sup> Marcel Müller,<sup>38</sup> Bozhao Nan,<sup>29</sup> Duc Nguyen,<sup>15</sup> Daniele Ongari,<sup>39</sup> Anthony Onwuli,<sup>16</sup> Can Özkan,<sup>40</sup> Sergio Pablo-Garcia,<sup>21</sup> Elton Pan,<sup>5</sup> Ratish Panda,<sup>41</sup> Sean Park,<sup>21</sup> Jaehee Park,<sup>18</sup> Dieter Plessers,<sup>42</sup> Tobias Plötz,<sup>24</sup> Ella M. Rajaonson,<sup>21</sup> Bojana Ranković,<sup>19</sup> Jarett Ren,<sup>43</sup> Rim Rihana,<sup>24</sup> Jurgis Ruza,<sup>5</sup> Akhil S. Nair,<sup>44</sup> Carter Salbego,<sup>22</sup> Erick Lopez Saldivar,<sup>45</sup> Arifin San,<sup>46</sup> Christina Schenk,<sup>47</sup> Stefan P. Schmid,<sup>48</sup> Dylan Schubert,<sup>49</sup> Philippe Schwaller,<sup>19</sup> Cher-Tian Ser,<sup>21</sup> Maitreyee Sharma Priyadarshini,<sup>50</sup> Yuxin Shen,<sup>51</sup> Kevin Shen,<sup>52</sup> Farshud Sorourifar,<sup>20</sup> Adrian Šošić,<sup>24</sup> Taylor Sparks,<sup>22</sup> Jan Christopher Spies,<sup>53</sup> Felix Strieth-Kalthoff,<sup>54</sup> Suraj Sudhakar,<sup>18</sup> Aditya Sundar,<sup>55</sup> Alessio Tamburro,<sup>56</sup> Clara Tamura,<sup>57</sup> Yifeng (Oliver) Tang,<sup>18</sup> Nikhil Thota,<sup>43</sup> Mohammad Erfan Toloue Sadegh Azadi,<sup>6</sup> Gary Tom,<sup>21</sup> Sang Truong,<sup>15</sup> Ricardo Valencia Alborno,<sup>51</sup> Luis Walter,<sup>58</sup> Lawrence Wang,<sup>59</sup> Fanjin Wang,<sup>60</sup> Andrew Wang,<sup>21</sup> Yiran Wang,<sup>5</sup> Jeffrey Watchorn,<sup>2</sup> Benjamin Weiser,<sup>6</sup> Geemi Wellawatte,<sup>19</sup> Alexander Wiecezorek,<sup>61</sup> Tim Würger,<sup>62</sup> Ilya Yakavets,<sup>2</sup> Jakob Zeitler,<sup>49</sup> Sylvester Zhang,<sup>6</sup> Yimu Zhao,<sup>2</sup> Yanqiao Zhu,<sup>32</sup> Ruijie (Ray) Zhu,<sup>18</sup> and Yunheng Zou<sup>63</sup>

<sup>1</sup>Department of Mechanical Engineering, Brigham Young University,  
701 E University Parkway, 350 EB, Provo, UT 84602, United States.

<sup>2</sup>Acceleration Consortium, University of Toronto,  
80 St George St, Toronto, ON M5S 3H6, Canada.\*

<sup>3</sup>Department of Manufacturing Engineering, Brigham Young University,  
265 Crabtree Technology Building (CTB), Provo, UT 84602, United States.

<sup>4</sup>ICodeGuru, Canal Road, Quaid-i-Azam Campus, Lahore, Punjab, Pakistan.

<sup>5</sup>Massachusetts Institute of Technology, Massachusetts Institute of Technology,  
77 Massachusetts Ave, Cambridge, MA 02139, United States.

<sup>6</sup>McGill University, 845 Rue Sherbrooke O, Montréal, QC H3A 0G4, Canada.

<sup>7</sup>Department of Computer Science, University College London,  
Gower Street, London WC1E 6BT, United Kingdom.

<sup>8</sup>Department of Statistical Science, University College London,  
Gower Street, London WC1E 6BT, United Kingdom.

<sup>9</sup>University of Melbourne, CSIRO, Parkville VIC 3052, Australia.

<sup>10</sup>Department of Energy Conversion and Storage, Technical University of Denmark,  
Anker Engelunds Vej 411, 2800 Kongens Lyngby, Denmark.

<sup>11</sup>Department of Chemistry, University of Oxford,  
Mansfield Rd, Oxford OX1 3TA, United Kingdom.

<sup>12</sup>National Research Council of Canada, 1200 Montreal Road, Ottawa ON, K1A 0R6, Canada.

<sup>13</sup>Ghana National Gas Company, Ghana National Gas Company,  
Gas House, 7th Avenue, Ridge West, Accra, Ghana.

<sup>14</sup>California Institute of Technology, 315 Annenberg IST Center,  
330 S Chester Ave, Pasadena, CA 91125, United States.

<sup>15</sup>Stanford University, 353 Jane Stanford Way, Stanford, CA, 94305, United States.

<sup>16</sup>Imperial College London, Exhibition Rd, South Kensington, London SW7 2AZ, United Kingdom.

<sup>17</sup>Atomics AI solutions private limited, Hyderabad, India.

<sup>18</sup>University of Chicago, 5801 S Ellis Ave, Chicago, IL 60637, United States.

<sup>19</sup>ETH Zurich, EPFL, Station 06, CH-1015 Lausanne, Switzerland.

<sup>20</sup>The Ohio State University, 281 W Lane Ave, Columbus, OH 43210, United States.

<sup>21</sup>Department of Chemistry, University of Toronto,  
27 King's College Circle, Toronto, ON M5S 1A1, Canada.

<sup>22</sup>University of Utah, University of Utah, 201 Presidents Circle, Salt Lake City, UT 84112, United States.

<sup>23</sup>Institute for Microelectronics and Microsystems (IMM),  
National Research Council (CNR), Institute for Microelectronics and Microsystems (IMM),

- National Research Council (CNR), Via del Fosso del Cavaliere 100, 00133 Rome, Italy.
- <sup>24</sup>Merck KGaA, Darmstadt, Germany, Frankfurter Straße 250, 64293 Darmstadt, Germany.
- <sup>25</sup>Princeton University, Department of Chemical and Biological Engineering, Princeton University, Princeton, New Jersey 08540, United States.
- <sup>26</sup>University of Queensland, Building 75, Cooper Rd, St Lucia QLD 4067, Australia.
- <sup>27</sup>Max Planck Institute for Dynamics of Complex Technical Systems, Sandtorstraße 1, 39112 Magdeburg, Germany.
- <sup>28</sup>FutureHouse, 1405 Minnesota St, San Francisco, CA 94107, United States.
- <sup>29</sup>University of Notre Dame, Holy Cross Dr, Notre Dame, IN 46556, United States.
- <sup>30</sup>Lawrence Berkeley National Laboratory, 1 Cyclotron Rd, Berkeley, CA 94720.
- <sup>31</sup>University Of Harvard, Harvard University, Massachusetts Hall, Cambridge, MA 02138, United States.
- <sup>32</sup>University of California, Los Angeles, 607 Charles E Young Dr E, Los Angeles, CA 90093, United States.
- <sup>33</sup>Department of Medical Physics & Biomedical Engineering, University College London, Gower Street, London WC1E 6BT, United Kingdom.
- <sup>34</sup>KAIST, KAIST (Korea Advanced Institute of Science and Technology), 291 Daehak-ro, Yuseong-gu, Daejeon 34141, South Korea.
- <sup>35</sup>University of Pittsburgh, 3700 O'Hara St, Pittsburgh, PA 15261, United States.
- <sup>36</sup>Cornell University, Ithaca, NY 14850, United States
- <sup>37</sup>Technical University of Denmark (DTU), Argous 64 & Potamou 77, Kifissia, Athens, Greece.
- <sup>38</sup>Merck KGaA, Darmstadt, Germany / University of Bonn, Frankfurter Straße 250, 64293 Darmstadt, Germany.
- <sup>39</sup>Syensqo, Viale Lombardia 20, Bollate, Italy.
- <sup>40</sup>TU Delft, Mekelweg 5, 2628CD Delft, the Netherlands.
- <sup>41</sup>Birla Institute of Technology, Mesra, Jharkhand 835215, India.
- <sup>42</sup>KU Leuven, Celestijnenlaan 200F, 3001 Leuven, Belgium.
- <sup>43</sup>Johns Hopkins University, 3400 N. Charles Street, Baltimore, MD 21218, United States.
- <sup>44</sup>Fritz-Haber Institute of the Max-Planck Society, Faradayweg 4-6, D-14195 Berlin, Germany.
- <sup>45</sup>Autonomous University of Chihuahua, Autonomous University of Chihuahua, Av. Escorza 900, Colonia Centro, 31000 Chihuahua, Chih., Mexico.
- <sup>46</sup>JSR Corporation, JSR Bioscience and informatics R&D center (JSR BiRD), 3-103-9 Tonomachi, Kawasaki-ku, Kawasaki, Kanagawa 210-0821, Japan.
- <sup>47</sup>IMDEA Materials Institute, Calle Eric Kandel 2, Tecnogetafe, 28906 Getafe (Madrid), Spain.
- <sup>48</sup>ETH Zurich, Vladimir-Prelog-Weg 1, 8093 Zurich, Switzerland.
- <sup>49</sup>Matterhorn Studio, Walton Street, Oxford, United Kingdom.
- <sup>50</sup>Virginia Tech, Virginia Tech, 925 Prices Fork Road, Blacksburg, VA 24061, United States.
- <sup>51</sup>University of Edinburgh, CH Waddington Building, King's Buildings, Max Born Crescent, Edinburgh, EH9 3BF, United Kingdom.
- <sup>52</sup>NobleAI, 345 California St Suite 600, San Francisco, CA 94104, United States.
- <sup>53</sup>University of Muenster, Corrensstraße 40, 48149 Münster, Germany.
- <sup>54</sup>University of Wuppertal, Germany, Gaußstraße 20, 42119 Wuppertal, Germany.
- <sup>55</sup>Argonne National Laboratory, Lemont, IL 60439, United States
- <sup>56</sup>Syensqo, 350 George Patterson blvd, Bristol, PA 19007, United States.
- <sup>57</sup>University of Washington, 3900 E Stevens Way NE, Seattle, WA 98195, United States.
- <sup>58</sup>Heidelberg University, Philosophenweg 19, 69120 Heidelberg.
- <sup>59</sup>Erthos, 2735 Matheson Blvd E 5, Mississauga, ON L4W 4M8, Canada.
- <sup>60</sup>Department of Mechanical Engineering, University College London, London, WC1E 7JE, United Kingdom.
- <sup>61</sup>Empa, Swiss Federal Laboratories for Materials Science and Technology, Ueberlandstrasse 129, Dübendorf, CH-8600 Switzerland.
- <sup>62</sup>Helmholtz-Zentrum Hereon, Max-Planck-Str. 1, 21502 Geesthacht, Germany.
- <sup>63</sup>University of Waterloo, 200 University Avenue West, Waterloo, Ontario, N2L 3G1., Canada.

The Acceleration Consortium and Merck KGaA hosted a 2-day virtual hackathon on March 27-28, 2024, bringing together scientists to explore, collaborate, and innovate in the field of Bayesian optimization for the physical sciences. Participants were encouraged to select or develop Bayesian optimization algorithms, apply them to benchmarking tasks, design new benchmarks, create instructional tutorials, and describe real-world applications. With over 100 participants across 69 academic, industry, and government organizations located in 59 cities, 19 countries, and 4 continents, this was a global event. The outputs from this event, including developed algorithms, benchmarks, and tutorials, will serve as valuable resources for the research community, in addition to the new skills learned and connections formed. Released projects and general information are available at <https://ac-bo-hackathon.github.io/> and other locations linked from individual project pages. This event demonstrates the potential of community-driven research efforts to accelerate advances in Bayesian optimization in chemistry and materials science.




Figure 1. Optimization traces for traditional design of experiment (DoE) methods compared with BO, typically outperforms. BO uses a smart model to predict where to look next in an experiment to find the best results with few experiments [1].

## I. Introduction

Bayesian optimization (BO) has emerged as a powerful tool in optimizing complex and expensive-to-evaluate functions, often outperforming traditional search methods in a variety of scientific domains such as optimizing composition and processing parameters to maximize alloy yield strength or identifying synthesis pathways that maximize efficacy of HIV drugs (Figure 1). Hackathons help people to connect, gain skills, and flesh out new ideas. In the words of Michelle Duke, the “Hackathon Queen”:

A hackathon is a short competition where people work together in teams to solve problems and challenges by coming up with solutions and ideas.

The goal of the AC BO Hackathon was to leverage the expertise of a diverse, global community to advance the development and application of BO techniques for solving critical challenges in the physical sciences. The hackathon also aimed to foster collaboration and knowledge sharing among participants from different backgrounds, including academia, national laboratories, government agencies, and private industry. The event attracted 139 active participants from 39 teams, representing 52 academic institutions, 3 national labs, and 14 companies. Likewise, the participants were located in 59 cities, 19 countries, and 4 continents (Figure 2). A full list of projects, including links to the corresponding GitHub repositories, submission video, and social media post are provided in Table III.

## II. Hackathon Details and Setup

This hackathon was in large part inspired by Jablonka *et al.* [2] and is one of several materials-focused hackathons in the past [3, 4].

Participants were provided with various resources to prepare for the hackathon—this included: GitHub classroom assignments with automated feedback, application- and theory-focused videos and tutorials, Python refresher materials, a list of tools to consider using during the hackathon, project submission instructions, and authorship criteria (Figure 3).

One of the unique aspects of this event is that it was hosted in Gather Town, a sort of union between traditional video conferencing software and retro arcade-style avatars and virtual spaces (Figure 4). Participants create a custom avatar and maneuver in a two-dimensional space. The videos and audio of other participants appear and become audible when nearby, and fade out when far away, simulating an in-person experience. At the beginning of the hackathon, all participants gathered to listen to keynotes in realtime, which were broadcasted via YouTube live and embedded into the Gather Town space. The videos were then made available on the hackathon website, which collectively garnered approximately 1600 views within two months. After the keynotes, teams were assigned tables in breakout rooms, each with a whiteboard. Individual tables were assigned as “private spaces” which isolated the shared audio and video within each space. This had a number of advantages for collaboration within and across teams.

The hackathon concluded with a project showcase accompanied by crowdsourced judging within a “poster room” (Figure 5).

Community judging occurred via Gavel, an automated pairwise comparison judging system. Use of this system helped to improve fairness, scalability, and accuracy by having judges compare projects relative to each other rather than assigning subjective numerical scores (for example, “rate your pain on a scale from 1 to 10, where 10 is the worst possible pain you can imagine”). Rather than scoring against a numeric rubric, judges answered a single holistic question for each assigned pair, effectively “which of these two projects is better?” The approach of pairwise judging reduces bias, allows for handling large competitions efficiently, and produces high-quality rankings using statistical models with dynamic assignments to judges to maximize information gain. It has been successfully used at HackMIT and other events to streamline judging and enhance transparency and credibility. See Table I for rankings of the top 10 projects and the awarded prizes.

The corresponding Gavel web app was hosted on Heroku according to directions in the Gavel repository. Individualized links were distributed to judges via email using Gavel’s SendGrid integration. Collectively, 35 judges cast 319 votes.

Table I. Ranked ~~Projects~~ [projects](#) with ~~Team Names~~ [team names](#) and ~~Prize Distribution~~ [prize distribution](#). ~~To~~ [The first prize value is the amount awarded to each team member and the value in parentheses is the maximum total per team: to avoid incentivizing single-person teams and very large teams, both per-person and per-team limits were imposed \(e.g., teams of 4 or more would have the max per-team amount divided equally rather than receive the max per-person amount\).](#)

| Rank | Proj. # | Team Name                 | Project Name                                                                                                                                                                            | <del>Prize*</del> <a href="#">Prize</a><br>(CAD) |
|------|---------|---------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------|
| 1st  | #23     | Noisy Nerds               | Reliable Surrogate Models of Noisy Data                                                                                                                                                 | 300 (1000 max)                                   |
| 2nd  | #34     | BOMS Prob                 | Streamlining Material Discovery - Bayesian Optimization in Thermal Fluid Mixtures                                                                                                       | 150 (500 max)                                    |
| 3rd  | #7      | Surface Science Syndicate | BayBE One More Time - Exploring Corrosion Inhibitors for Materials Design                                                                                                               | 75 (250 max)                                     |
| 4th  | #5      | KLM                       | Comparing Bayesian Optimization Methods ... Against Simulated "Human" Decision-making                                                                                                   | 40 (125 max)                                     |
| 5th  | #8      | Molecular Representation  | BO for Drug Discovery - What is the role of molecular representation?                                                                                                                   | 40 (125 max)                                     |
| 6th  | #9      | PME No Hikari             | Optimizing The CO2 Adsorption Capacity of Metal-Organic Frameworks Using Thompson Sampling                                                                                              | 40 (125 max)                                     |
| 7th  | #11     | BlendDS                   | BlendDS - An intuitive specification of the design space for blends of components                                                                                                       | 40 (125 max)                                     |
| 8th  | #30     | SERO Opt                  | Active learning for voltammetry waveform design                                                                                                                                         | 40 (125 max)                                     |
| 9th  | #43     | General Optimizers        | Bayesian Optimization for Generality                                                                                                                                                    | 40 (125 max)                                     |
| 10th | #3      | Sparks Group              | Take Your Time - <del>Measuring Optimization Performance as a Function of</del> <a href="#">Improving Optimization Performance Through Greater Investment in</a> ACQF Optimizer Runtime | 40 (125 max)                                     |

Table II. Project Topics for the Hackathon, [with the number of submissions whose project page records each topic](#). [Five project pages did not record a topic](#). See [the submission page](#) for more details.

| — Topic                     | Description                                                                             | <a href="#">Projects</a> |
|-----------------------------|-----------------------------------------------------------------------------------------|--------------------------|
| 1 <b>Apply Algorithms</b>   | Choose an algorithm and apply it to a <a href="#">hackathon benchmark task</a>          | <a href="#">5</a>        |
| 2 <b>Develop Benchmarks</b> | Develop a new benchmark and add it to a suite of benchmarks                             | <a href="#">2</a>        |
| 3 <b>Create Tutorials</b>   | Create "gentle introduction" tutorials for <a href="#">advanced optimization topics</a> | <a href="#">4</a>        |
| 4 <b>Propose Tasks</b>      | Propose materials tasks that <i>can</i> and <i>should</i> be tackled with BO            | <a href="#">9</a>        |
| 5 <b>General</b>            | Other projects related to BO for the physical sciences                                  | <a href="#">20</a>       |




Figure 2. Demographic distributions of the hackathon participants and their affiliations. The hackathon supported 139 participants across 69 academic, industry, and government organizations located in 59 cities, 19 countries, and 4 continents.

**Hackathon resources**

- Slack - Join the hackathon's Slack channel to connect, ask questions, and form teams
- Project proposals - team leaders should submit their project proposals using the instructions here
- Gather - Add the Gather Town event to your calendar
- Accelerated discovery forum - create an account to join the discussion forums
- Bayesian optimization - Watch the Accelerate '23 training workshop on Bayesian optimization below

**Getting Started**

**Orientation Modules**

Please complete the following orientation assignments in preparation for the hackathon to familiarize the tools and concepts you'll need. You will need to create a GitHub account to access these resources. If you are already familiar with git and GitHub, you may skip directly to "Intro to GitHub Classroom". When prompted, please select an identifier from the list that's given. In addition to these orientation modules, we also recommend that you familiarize yourself with Markdown syntax if this is new to you.

- Intro to Git and GitHub
- Intro to GitHub Classroom
- Python Refresher
- Simple Bayesian optimization script

**Bayesian optimization**

Bayes Rule by 3Blue1Brown - explanation of the fundamental concept behind Bayesian methods

This is Bayes' rule

A Visual Exploration of Gaussian Processes - An explanation of a commonly used predictive model in Bayesian optimization with visualizations and mathematical rigor

Exploring Bayesian Optimization - An overview of active learning and Bayesian optimization with visualizations and mathematical rigor

Bayesian Optimization Book - For those that REALLY want to understand BO, this is a self-contained, comprehensive deep dive that builds up basic principles from "scratch" and delves into advanced topics. The book is free and open-source.

**Python**

Programming expertise is not required, but at least beginner Python programming experience is recommended for participation in code focused projects of the hackathon. For those looking for a brief, interactive refresher on Python programming, see the GitHub Classroom assignment from the first section on this page. For those without prior Python experience, we recommend you complete an introductory Python course in preparation for the hackathon. Some resources are listed below.

|                                                     |                                                                      |
|-----------------------------------------------------|----------------------------------------------------------------------|
| Python for Everybody by University of Michigan      | Learn Python 3 by CodeAcademy                                        |
| Intro to Python Programming for Materials Engineers | Python Beginners Guide for Programmers by Python Software Foundation |

If you have no prior programming experience, you may wish to start with the Python Beginners Guide for Non-programmers by Python Software Foundation.

**BO Tools**

Use of the tools listed on this page is not a requirement. A diverse set of packages and implementations is encouraged. Likewise, multiple teams using the same package is not a problem, in part because implementations can remain private during the course of the hackathon.

If you'd like to see a specific tool listed here, please navigate to the "improve this page" link at the bottom of the page and open a pull request. See also the Acceleration Consortium's curated list of optimization tools.

**BayBE**

Bayesian Back End

This hackathon is co-organized by Merck KGaA. You may consider using an internal Merck tool that was recently open sourced, called the Bayesian optimization Backend (BayBE). This tool focuses on real world experimental campaigns in the physical sciences with support for integrating domain knowledge through custom parameter encodings and built in chemical encodings. In addition, it allows for custom surrogate models, mixed variable types, and transfer learning.

Figure 3. A snapshot of resources listed on the hackathon webpage such as hackathon orientation, intro to BO, and a Python refresher assignment. These resources prepared participants to maximize their time during the two-day synchronous portion of the hackathon and helped level the playing field for participants with varied backgrounds and skill levels. Image is Portions of the screenshot have been blurred per preprint server policy to protect potentially identifying information.




Figure 4. Gather town [keynote room](#) (left), custom avatars (top-right), and an example of a breakout room for teams (bottom-right). Keynotes were broadcasted in realtime to participants via an embedded YouTube livestream. Use of Gather Town helped level the playing field for teams who were in physically separate locations and made it easier for facilitators and other teams to have more natural “check-ins” with other projects. [Some participant display names have been pixelated to protect potentially identifying information.](#)



Figure 5. The synchronous portion of the hackathon concluded with a poster session and community judging. One participant noted that “it almost felt like a real poster session.” [Participant display names have been pixelated to protect potentially identifying information.](#)

Table III: List of projects with links to GitHub, Social Media, and Video.  
Project pages on the hackathon website are available at <https://ac-bo-hackathon.github.io/projects/>.

| <u>Proj. #</u> | <u>Project Name</u>                                                                                                                                | <u>Links</u>                                                                                                                                                                                                                                                      |
|----------------|----------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1              | Multi-objective Benchmarking of Dragonfly against BoTorch                                                                                          |          |
| 2              | Long-run Behaviour of Multi-fidelity Bayesian Optimisation                                                                                         |          |
| 3              | Take Your Time - Improving Optimization Performance Through Greater Investment in ACQF Optimizer Runtime                                           |                                                                                             |
| 4              | SimpleGPT-BO, Simplified GPT-Powered Bayesian Optimization                                                                                         |          |
| 5              | Comparing Bayesian Optimization Methods Across Multiple Hyperparameters Against Simulated "Human" Decision-making                                  |                                                                                             |
| 6              | Multi-Objective Bayesian Optimization for Transparent Electromagnetic Interference Shielding with Thin-Film Structures                             |          |
| 7              | BayBE One More Time - Exploring Corrosion Inhibitors for Materials Design                                                                          |          |
| 8              | BO for Drug Discovery-What is the role of molecular representation?                                                                                |          |
| 9              | Optimal MOF Selection for CO2 capture using Thompson sampling                                                                                      |                                                                                             |
| 10             | Navigating the black box of zeolite synthesis with Bayesian Optimization                                                                           |          |
| 11             | BlendDS - An intuitive specification of the design space for blends of components                                                                  |          |
| 12             | Robust GPs for Sustainable Concrete via Bayesian Optimization                                                                                      |          |
| 13             | Interpretability of Bayesian Optimisation Campaigns                                                                                                |                                                                                             |
| <u>14</u>      | <u><a href="#">Bayesian optimization of likely negative candidates in imbalanced biological datasets</a></u>                                       |                                                                                                                                                                                |
| 15             | Adaptive Batch Sizes for Bayesian Optimization of Reaction Yield                                                                                   |                                                                                             |
| 16             | BOPE-GPT, Preference Exploration with the curious AI chemist                                                                                       |          |
| 17             | Comparative Analysis of Acquisition Functions in Bayesian Optimization for Drug Discovery                                                          |       |
| 18             | Investigation of Multi-Objective Bayesian Optimization of QM9 Dataset                                                                              |    |
| <u>19</u>      | <u><a href="#">Quantum Bayesian Optimization for Automatic Chemical Design</a></u>                                                                 |                                                                                                                                                                                                                                                                   |
| 20             | Closed loop optimization of hydrogel formulations using dynamic light scattering                                                                   |    |
| 21             | Benchmarking Molecular Descriptors with Actively Identified Subsets (MolDAIS)                                                                      |                                                                                         |
| 22             | Chemical Similarity-Informed Earth Mover's Distance Kernel Bayesian Optimization for Predicting the Properties of Molecules and Molecular Mixtures |                                                                                         |
| 23             | Reliable Surrogate Models of Noisy Data                                                                                                            |                                                                                                                                                                              |
| 24             | ScattBO Benchmark - Bayesian optimisation for materials discovery                                                                                  |    |
| 25             | Bayesian Optimized De Novo Drug Design for Selective Kinase Targeting                                                                              |    |
| 26             | Multiple-Context Bayesian Optimization                                                                                                             |    |
| 27             | How does initial warm-up data influence Bayesian optimization in low-data experimental settings?                                                   |    |
| 28             | The Impact of Dataset Size on Bayesian Optimization, Insights from the QM9 Dataset                                                                 |                                                                                         |
| <u>29</u>      | <u><a href="#">A Bayesian Approach to Predict Solubility Parameters</a></u>                                                                        |                                                                                                                                                                                                                                                                   |
| 30             | Active learning for voltammetry waveform design                                                                                                    |    |
| 31             | A tutorial on ask/tell mode for Ax                                                                                                                 |    |
| 32             | Efficient Protein Mutagenesis using Bayesian Optimization                                                                                          |    |
| 33             | Bayesian Optimization for Hyperspectral Co-heritability Search                                                                                     |    |
| 34             | Streamlining Material Discovery - Bayesian Optimization in Thermal Fluid Mixtures                                                                  |                                                                                         |
| 35             | Tutorial for GAUCHE - A Library for Gaussian Processes in Chemistry                                                                                |    |
| 36             | Scalable Nonmyopic Bayesian Optimization in Dynamic Cost Settings                                                                                  |    |
| 37             | The Effects of Post-Modelling Performance Metric Computation on the Efficiency of Bayesian Optimizers                                              |    |

*Continued on the next page*

Table III, continued.

| Proj. # | Project Name                                                                                                         | Links                                                                                                                                                                                                                                                       |
|---------|----------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 38      | Bayesian methods in symbolic regression                                                                              |                                                                                       |
| 39      | Divide and Conquer - Local Gaussian Processes to design Covalent Organic Frameworks for Methane Deliverable Capacity |    |
| 40      | Optimizing Chemical Reaction Conditions with Multi-Agent Systems Using LLM and BO                                    |    |
| 41      | Retrieval Augmented Bayesian Optimization                                                                            |    |
| 42      | <a href="#">Bayesian Optimization and Design of Experiments for Pharmaceutical Formulation (BODoE)</a>               |                                                                                                                                                                                                                                                             |
| 43      | Bayesian Optimization Awesome List                                                                                   |                                                                                       |
| 44      | Rank-based Bayesian Optimization                                                                                     |    |
| 45      | Bayesian Optimization for generality                                                                                 |                                                                                       |

### III. Projects’ Key Findings

This section provides a ~~comprehensive summary and highlights summary of~~ the key findings from all 45 project submissions. To streamline the ~~evaluation process, all YouTube video submissions were transcribed and analyzed using process, the speech transcript of each YouTube video submission was retrieved and condensed into a draft summary by~~ an AI agent powered by ~~Anthropies’~~ Anthropic’s *claude-3-5-sonnet-20240620* language model with a temperature of 0.3, ~~ensuring consistent and accurate processing. Each transcript was then manually edited to ensure a high-quality manuscript compilation. This automated approach enhances efficiency while maintaining a structured and objective assessment of the submissions~~ applied uniformly across submissions. Each draft was then reviewed and edited by the organizers to correct errors and ensure a faithful description; the per-project summaries are archived alongside the manuscript source. Six projects (14, 19, 23, 29, 34, and 42) had no recorded video submission to transcribe; their summaries were instead prepared from the corresponding project pages and each team’s public artifacts, as noted at the end of each such summary.

#### Project 1: Multi-objective Benchmarking of Dragonfly against BoTorch

This research project compared the performance of two BO packages, Dragonfly [5] and BoTorch [6], in a multi-objective optimization context. The researchers optimized the Branin-Currin test function with added noise, conducting 20 trials of BO over different random seeds. Each trial used a batch size of 1 and consisted of 20 iterations, evaluating 20 candidates per iteration. The results showed that BoTorch’s acquisition functions, particularly the q-noisy expected hypervolume improvement (qNEHVI) [7], outperformed Dragonfly’s acquisition function. The qNEHVI acquisition function achieved the lowest log hypervolume difference on average across the trials, while Dragonfly’s acquisition function struggled to surpass a

randomly generated set of candidates. Visualization of the average Pareto front produced by different acquisition functions further demonstrated BoTorch’s superiority, with its acquisition functions covering a larger portion of the Pareto front compared to Dragonfly. The researchers concluded that BoTorch outperformed Dragonfly within the scope of their study, but noted the need for further investigation using different batch sizes, input and output constraints, and test functions simulating chemical reactions to apply the findings to material design problems.

#### Project 2: Long-run Behaviour of Multi-fidelity Bayesian Optimisation

This research project investigates the long-term performance of multifidelity Bayesian optimization (MFBO) methods [8] compared to single-fidelity Bayesian optimization (SFBO) methods. The study focuses on scenarios where two fidelities are available: an accurate but expensive target fidelity that fully represents the objective, and a less accurate but cheaper low fidelity. The researchers observed that while MFBO methods initially outperform SFBO up to a budget of approximately 20, they begin to underperform beyond this point. This behavior is particularly evident with linear kernel MFBO methods, where the performance gap becomes more pronounced after a cost of around 60. The project aims to evaluate the factors contributing to the long-term issues of MFBO methods. Specifically, the researchers plan to investigate the impact of kernel choice, frequency of low fidelity queries, and acquisition function selection on MFBO performance. By analyzing these aspects, the study seeks to identify potential improvements for MFBO methods to enhance their applicability in real-life optimization tasks. The ultimate goal is to develop MFBO approaches that can maintain their performance advantage over SFBO methods in the long run, leveraging the cost-effectiveness of low fidelity queries while achieving superior optimization results.

### **Project 3: Take Your Time - Improving Optimization Performance Through Greater Investment in ACQF Optimizer Runtime**

This research project explores an enhancement to standard BO campaigns by introducing a random seed approach to improve optimization outcomes. The team from the University of Utah’s Material Science and Engineering department investigated the variability in BO campaigns when starting from the same initial data points. The key innovation in this work is a custom approach to calculating the acquisition function. Instead of computing it once per iteration, as is standard in most BO methods [9, 10], the researchers calculate the acquisition function multiple times using different random seeds. They then select the highest-performing result at each iteration. This method, termed “random retries Optimizer”, showed significant improvements in optimization performance, particularly in avoiding local optima. The researchers demonstrated that this approach consistently achieves near-optimal performance among the possible campaigns stemming from the initial data points. While this method requires additional computational time for the multiple acquisition function calculations, the performance gains appear to outweigh the increased computational cost. The team is continuing to investigate questions such as the trade-off between additional compute time and performance improvements, as well as how the method performs across different types of optimization problems, including simpler ones compared to the challenging example presented.

### **Project 4: SimpleGPT-BO, Simplified GPT-Powered Bayesian Optimization**

SimpleGPT-BO is a ChatGPT-integrated tool that makes Bayesian Optimization accessible without installation requirements. The system provides guided optimization through multiple acquisition functions (PI, EI, UCB, TPE) and supports multi-objective optimization via weighted single objective transformation. Users access the tool directly through ChatGPT’s interface or via provided URLs, benefiting from pre-generated instructions and global availability. While designed for straightforward optimization tasks rather than complex analyses, the open-source project (MIT License) welcomes community contributions to expand functionality and applications. The tool aims to democratize Bayesian Optimization by eliminating technical barriers and providing step-by-step guidance for users across all experience levels.

### **Project 5: Comparing Bayesian Optimization Methods Across Multiple Hyperparameters Against Simulated “Human” Decision-making**

This research project investigated the effectiveness of BO methods that mimic human experimentalist approaches in materials science [11], comparing them to traditional BO techniques and random search baselines [10, 12]. The study focused on three key axes:

model complexity, number of features, and acquisition function. The researchers found that a simplified “human experimentalist” approach, using automated feature engineering with a linear model, three features, and an exploitative acquisition function, was competitive with more complex random forest models in terms of enhancement factor and acceleration factor. Both the simplified approach and random forest models outperformed random search. Interestingly, for linear models, increasing complexity through additional features or more exploratory acquisition functions led to decreased performance. The main conclusion was that low-dimensional, interpretable models were comparable to traditional BO methods for the datasets and hyperparameter regimes studied. The researchers suggest future work could involve comparing their results to actual human performance and improving their methodology using theoretical approaches.

### **Project 6: Multi-Objective Bayesian Optimization for Transparent Electromagnetic Interference Shielding with Thin-Film Structures**

This research project focuses on applying multi-objective BO to design transparent electromagnetic interference (EMI) shielding using thin film structures [13]. The primary objectives are to maximize both transmittance and shielding effectiveness simultaneously. The optimization problem involves selecting materials and thicknesses for each layer of the thin film structure, with a search space comprising 12 material choices (Ag, Al, Al<sub>2</sub>O<sub>3</sub>, Cr, Ni, Pd, Si<sub>3</sub>N<sub>4</sub>, SiO<sub>2</sub>, Ti, TiN, TiO<sub>2</sub>, W) and thickness ranges from 5 to 20 nm. No pre-existing dataset was involved: both objectives were computed on demand with LayerLumos, the team’s open-source transfer-matrix simulator for the optical and radio-frequency response of candidate film stacks. The researchers implemented multi-objective BO using random scalarization for the acquisition function via BayesO package [14]. Gaussian process regression and expected improvement were employed as part of the BO framework. The results demonstrate that the multi-objective BO approach effectively identified the Pareto frontier for the two objectives of transmittance and shielding effectiveness. This suggests that the method is successful in finding optimal trade-offs between transparency and EMI shielding performance for thin film structures, which could have applications in areas such as spacecraft windows where both properties are crucial.

### **Project 7: BayBE One More Time - Exploring Corrosion Inhibitors for Materials Design**

This research project explored the use of BO for efficiently identifying effective corrosion inhibitors in materials design. The researchers utilized the BO implementation from the BayBE package [15] to evaluate its performance on multiple experimental datasets [16] involving different alloys, primarily aluminum-based. The study found that the choice of molecular descriptor en-

coding significantly impacted the optimization performance, especially for larger datasets like A1000 and A2024. The SMILES-based approach with Mordred encoding [17] outperformed the other methods for these datasets, while random sampling underperformed. Interestingly, for the A775 dataset, which had a sparse distribution of reported efficiencies and fewer data points, this trend was not observed. The researchers also explored transfer learning capabilities, using information from the A2024 alloy dataset to inform the optimization process for the A1000 dataset. This transfer learning approach outperformed the uninformed method after only 15 iterations, demonstrating its potential for accelerating the discovery of effective corrosion inhibitors in experimental settings.

#### **Project 8: BO for Drug Discovery-What is the role of molecular representation?**

This research project investigated the impact of molecular featurization methods on BO performance for guiding molecular experiments. The team explored alternatives to the commonly used connectivity-based fingerprints, aiming to determine which featurization techniques yield the best results in BO for molecular discovery. The study found that without modifying the features, MACCS [18] and RDKit [19] featurization methods outperformed the default options presented in previous publications. To address the high dimensionality of molecular features, the researchers employed specialized Gaussian processes and explored random forest surrogates. However, the random forest approach did not improve performance and required more computational resources. Principal Component Analysis (PCA) was successfully used to enable Gaussian processes to handle large molecular featurizations. Notably, physicochemical featurizations like RDKit and Mordred [17] outperformed the previous benchmark winner, MolDQN [20]. The team’s findings were consistent with benchmarks in the BO library but suggested that RDKit could produce even better results than Mordred. Overall, the research highlights the importance of careful featurization selection in BO for molecular discovery, with RDKit emerging as a particularly effective option.

#### **Project 9: Optimal MOF Selection for CO<sub>2</sub> capture using Thompson sampling**

This research project focuses on optimizing the selection of Metal-Organic Frameworks (MOFs) for carbon capture applications using BO, specifically Thompson Sampling. MOFs are nanoporous materials with high potential for carbon capture [21, 22], but their synthesis and characterization are expensive and time-consuming [23]. The goal was to develop an efficient method for identifying high-performing MOF candidates using a small dataset. The team employed a Gaussian Process model trained on the CRAFTED dataset [24], using both RACs features and geometric features of MOFs, with CO<sub>2</sub> uptake as the output. Thompson Sampling was used as

the acquisition function to select the next best sample based on the posterior distribution, inherently performing a Bayesian update. The results showed that Thompson Sampling was twice as efficient as random sampling in identifying high-performing MOF candidates. This approach is notable for being the first application of Thompson Sampling to MOF candidate selection, and it requires no hyperparameter tuning, making it easily transferable and cost-efficient. The potential impact of this method is significant, as it could accelerate MOF design for direct air capture applications, with a goal of absorbing 30 million tons of CO<sub>2</sub>, and could also be applied to post-combustion capture scenarios.

#### **Project 10: Navigating the black box of zeolite synthesis with Bayesian Optimization**

This research project explores the application of BO to zeolite synthesis, a complex process with significant industrial importance. Zeolites are crystalline materials composed of interconnected silicate and aluminate tetrahedra, widely used as absorbents and catalysts [25]. The synthesis of zeolites involves multiple parameters, including silicon and aluminum sources, organic molecules, water, temperature, and time. The process aims to achieve specific properties such as high crystallinity, large external surface area, and particular silicon-to-aluminum ratios, while also considering economic factors like synthesis temperature, duration, and ingredient concentrations [26]. The project presents a GitHub repository containing an introductory text on zeolites and their synthesis, along with a Jupyter notebook demonstrating BO applications. The notebook is divided into two main sections: the first optimizes an analytical dummy function using zeolite synthesis parameters, exploring various aspects of BO including continuous, categorical, and mixed variable types, parameter constraints, and single and multiple objectives. The second section applies BO to propose new experiments based on existing literature data. This approach aims to accelerate the traditionally time-consuming trial-and-error process of zeolite synthesis optimization, potentially reducing costs and improving efficiency in industrial applications.

#### **Project 11: BlendDS - An intuitive specification of the design space for blends of components**

This project focuses on developing an interface for BO that bridges the gap between domain experts (i.e. chemists and materials scientists) and machine learning algorithms. The system allows scientists to specify experimental parameters and constraints in natural language, which is then translated into a structured dictionary format by a large language model. This dictionary is subsequently converted into a Python object that can be used for design of experiments and optimization tasks. The key feature of this interface is its ability to generate diverse trial designs for efficient sampling of the experimental space. It incorporates dimensionality reduction techniques to visualize the design space, enabling

users to select the most diverse and informative trials. This approach aims to maximize the information gained from each experiment, potentially reducing the number of trials needed to reach optimal results. The project is open-source, inviting contributions and integration with existing BO frameworks.

#### **Project 12: Robust GPs for Sustainable Concrete via Bayesian Optimization**

This project introduces Robust Gaussian Processes (Robust GPs) [27] to the concrete sciences experiments with Bayesian optimization campaigns. The research addresses the common issue of outliers in concrete science experiments by applying Robust GPs to account for these anomalies. The project provides a tutorial and accompanying code in a GitHub repository, demonstrating how to implement Robust GPs in this specific context. The study utilizes a dataset provided by Meta [28], which includes input variables and a target variable representing concrete strength. Initial data analysis reveals potential outliers in the strength variable, with values between 2 and 2,000. The researchers then apply Robust GP to this dataset, achieving good predictive performance. The project culminates in demonstrating how to obtain Bayesian recommendations using the Robust GP as a surrogate model, providing a practical approach for optimizing concrete mixtures while accounting for experimental outliers.

#### **Project 13: Interpretability of Bayesian Optimisation Campaigns**

This research project focused on developing novel methods for interpreting Bayesian Optimization campaigns by incorporating the temporal component often neglected in end-of-campaign analyses. The study utilized data from a self-driving lab experiment aimed at optimizing the conductivity of a coating. The researchers recreated the Gaussian Process (GP) model employed in the original experiment, training it on subsets of data to simulate mid-campaign conditions. Two primary methods were investigated: First, cross-sections of the GP model were taken at various stages of the campaign, holding all but one feature constant. This approach revealed that closer nozzle distances produced better results, aligning with prior beliefs and serving as a potential mid-campaign sense check to validate model alignment with previous experimentation. The second method involved predicting the error of the next sample based on the model trained on data collected up to that point. This analysis showed that performance improvement plateaued around the time the optimal sample was produced, suggesting its potential use as an early stopping criterion if error is monitored throughout the campaign. The research demonstrates the value of incorporating temporal analysis in Bayesian Optimization campaigns, offering insights that could enhance decision-making during experiments. The cross-sectional approach provides a means to validate model behavior

against prior knowledge, while the error prediction method could inform stopping criteria, potentially improving efficiency in optimization processes. These techniques offer promising avenues for real-time interpretation and guidance in Bayesian Optimization experiments, particularly in self-driving lab contexts.

#### **Project 14: Bayesian optimization of likely negative candidates in imbalanced biological datasets**

High-throughput biological screens yield strongly class-imbalanced data because reliable negative examples are hard to obtain, which motivates positive-unlabeled (PU) learning [29]. This team asked whether BO can choose the reliable negatives, treating the labeling of unlabeled peptides as the decision variable: at each iteration the current classifier scores the unlabeled pool, an acquisition function defined over the predicted class-probability distribution (expected improvement or an upper confidence bound) ranks the candidates, the top-ranked sequences are committed as negatives, and the classifier is retrained, with held-out AUC as the objective. On the hemolysis task (1826 positive, 7490 negative and 7295 unlabeled sequences), a classifier trained on the labeled data alone reached AUC 0.71 to 0.73. With a random-forest surrogate, the BO-selected negatives raised AUC to 0.78 within two iterations and held it through the seventh, after which AUC decayed to between 0.67 and 0.74 as the committed negative set grew, while accuracy rose monotonically from 0.02 to 0.40 over 20 iterations. A neural-ensemble variant reproduced the early gain (AUC 0.65 to 0.70) but cost minutes per retraining, limiting iterations. A second target, binding to the SHP-2 protein, was set up to compare EI, UCB and random selection but completed only one iteration within the event. (No video was submitted; summary prepared from the team’s project page and repository notebooks.)

#### **Project 15: Adaptive Batch Sizes for Bayesian Optimization of Reaction Yield**

This research project investigates the optimization of batch sizes in BO for chemical reaction yield optimization. All optimization runs were retrospective queries of the 1,728-reaction palladium-catalyzed direct arylation benchmark of Shields et al. [30], which spans 12 ligands, 4 bases, 4 solvents, 3 concentrations, and 3 temperatures with measured yields; the project code fetches this dataset at run time from the EDBO+ repository [31] rather than shipping a copy of its own. The study focuses on determining the optimal batch size ( $Q$ ) at each step of the BO process, considering both smaller and larger batch sizes. The researchers examine how  $Q$  impacts the cost of retraining models, such as large language models, and how it relates to additional experimentation overhead. The project findings indicate that the optimal  $Q$  depends on a trade-off between model retraining time and additional experimentation overhead. For a fixed batch size, the researchers observed two minima at  $Q=3$  and

$Q=7$ , reflecting the relative importance of  $Q$  in different scenarios. With very large batch sizes, the retraining cost had minimal impact on the overall experimental overhead, while experimentation overhead had a significant influence. The study also explored adaptive batch sizing strategies, decreasing batch sizes when performance was good and increasing them as the surrogate model became more trustworthy. The researchers conclude that batch sizes have a crucial effect on BO efficiency in real-world settings, and adaptive batch sizes could effectively balance the trade-off between model retraining and batch sampling effectiveness.

#### **Project 16: BOPE-GPT, Preference Exploration with the curious AI chemist**

This research project focused on optimizing the Fischer-Tropsch synthesis process [32], which converts syngas into biofuel hydrocarbons. The optimization problem on the dataset [33, 34] involved four input variables (space time, syngas ratio, temperature, and pressure) and four output variables (carbon conversion, selectivity, methane to paraffins, and olefins). The researchers proposed a novel approach using preferential BO [35] with large language models (LLMs) to prioritize outputs based on stakeholder preferences. The methodology employed the typical BO process powered by BoTorch, using Expected Utility as the acquisition function. However, the innovation lay in using an LLM (specifically GPT-4) to perform pairwise comparisons for output prioritization. The results showed that outputs exhibited well-defined optima trade-offs and some degree of monotonicity, which was important for preferential BO. The researchers tested three objective scenarios: optimizing all outputs equally, maximizing only CO conversion, and maximizing three objectives while minimizing olefins. For the first two objectives, the LLM-based preferential BO performed similarly to traditional utility function-based methods. However, for the third, more complex objective, the LLM approach performed worse than the utility function method and only slightly better than random exploration, suggesting limitations in the LLM’s ability to handle complex optimization scenarios. The project also included the development of an interactive application for visualizing BO results and running live optimizations.

#### **Project 17: Comparative Analysis of Acquisition Functions in Bayesian Optimization for Drug Discovery**

This project investigates the application of BO techniques directly on molecular fingerprints for drug discovery, focusing on comparing different surrogate models, acquisition functions and small, diverse, unbalanced, and noisy datasets [36]. The researchers used the LD50 dataset from PTDC with SMILES encodings [37, 38], which were transformed into ECFP fingerprints [39], creating a 2048-dimensional feature space. Two surrogate models were examined: Gaussian Processes (GP) and

Random Forests (RF). The results showed that in this high-dimensional space, Random Forests with uncertainties evaluated as variance between different trees achieved expected performance across all acquisition functions, significantly outperforming random selection. In contrast, Gaussian Processes failed to perform well, likely due to the challenges posed by the high-dimensional space ( $2^{2048}$  possible combinations) and the relatively small dataset of only 7,000 data points. The researchers concluded that machine learning models, particularly Random Forests with various acquisition functions, perform well when dealing with high-dimensional molecular fingerprint spaces, demonstrating their potential for drug discovery applications with certain selection biases.

#### **Project 18: Investigation of Multi-Objective Bayesian Optimization of QM9 Dataset**

In this work, the team explores the application of multi-objective Bayesian optimization (MoBO) techniques to molecular property optimization using the QM9 dataset [40]. The project implements and compares state-of-the-art MoBO algorithms—including Expected Hypervolume Improvement (EHVI) [41] and ParEGO—to [42], [42], to efficiently navigate the trade-offs between multiple molecular objectives such as HOMO-LUMO gap, molecular weight, and dipole moment. Leveraging molecular feature encodings and Gaussian process surrogate models, the optimization framework identifies candidate molecules that balance competing objectives in chemical design. This study highlights the potential of MoBO to accelerate lead discovery by providing a principled, sample-efficient method for exploring high-dimensional molecular spaces.

#### **Project 19: Quantum Bayesian Optimization for Automatic Chemical Design**

This project proposed a quantum-computing route to BO for molecular design, with quantum Gaussian-process regression [43] in place of the classical surrogate. The submission was never carried out, so no dataset was used.

#### **Project 20: Closed loop optimization of hydrogel formulations using dynamic light scattering**

This research project proposes a self-driving lab approach to optimize hydrogel formulations using BO and automated characterization techniques. The study aims to address the challenge of creating hydrogels with specific properties for applications such as cell culture, tissue engineering, drug delivery, and agriculture. The proposed workflow utilizes an open-source liquid handling robot to mix hydrogel formulations in a 96-well plate, followed by cross-linking using various methods (LED array, heating module, or time). The key innovation is the use of Dynamic Light Scattering (DLS) [44] for automated characterization of gelation and viscoelastic properties, which can measure gel stiffness up to 10 kilopascals. This DLS method is implemented using a plate

reader specifically designed for 96-well plates, enabling high-throughput analysis. The resulting property data is then fed into a Bayesian Optimizer to determine the next set of formulations and processing parameters to test. This approach aims to efficiently explore the complex relationship between hydrogel formulation parameters and their resulting properties, potentially accelerating the development of custom hydrogels for specific applications in biological research and improving data availability and reliability in the field. As a two-day proposal, the closed loop was exercised only in software: the accompanying example script populates a BayBE [15] campaign with mock measurement values, and no experimental hydrogel dataset was generated or used.

#### **Project 21: Benchmarking Molecular Descriptors with Actively Identified Subsets (MolDAIS)**

This research presents a novel approach called ~~MOLDES~~-MolDAIS (Molecular Descriptors with Actively Identified ~~Subspaces~~Subsets) for molecular property optimization. The method addresses the challenge of optimizing molecules in high-dimensional spaces by using molecular descriptors - sets of rotationally and translationally invariant calculations performed on molecular graphs - coupled with active subspace identification. ~~MOLDES~~-MolDAIS employs a sparse axis-aligned subspace Gaussian Process prior, which actively learns an encoding while performing Bayesian optimization. Recent works [45, 46] are increasingly turning towards active encoding of molecular feature spaces. The researchers evaluated ~~MOLDES~~-MolDAIS on three case studies: experimental lipophilicity (~~measurements for 4,200 compounds~~ from the MoleculeNet benchmark collection [47]), log P optimization ~~benchmark~~ (over the 250,000-molecule ZINC subset popularized by Kusner et al. [48], and computed power conversion efficiencies for 30,000 molecules), and ~~power conversion efficiency organic photovoltaic candidates~~ from the Harvard Clean Energy Project (30,000 compounds) [49]. In all cases, ~~MOLDES~~-MolDAIS demonstrated superior performance compared to other optimizers, particularly in larger datasets. For the log P optimization, ~~MOLDES~~ MolDAIS consistently found the optimal molecule within 100 iterations. The method also showed strong performance in constrained optimization problems, often achieving the best-case scenario and maintaining a favorable worst-case scenario compared to other methods. Overall, ~~MOLDES~~-MolDAIS proved efficient in identifying high-performing molecules in low-data regimes, offering a promising approach for molecular property optimization tasks.

#### **Project 22: Chemical Similarity-Informed Earth Mover’s Distance Kernel Bayesian Optimization for Predicting the Properties of Molecules and Molecular Mixtures**

This research project focuses on developing chemical similarity-informed distance functions and kernels for

explainable Bayesian optimization, specifically targeting the prediction of properties for molecular mixtures. The researchers propose a novel approach that bypasses the need for embedding vectors by directly providing pairwise distances between data points in the kernel function of a Gaussian Process (GP) model [50]. The project introduces the Earth Mover’s Distance (EMD) kernel [51] into the GP framework to calculate pairwise distances between mixtures based on individual component distances. This method was tested for predicting yields of binary reactant mixtures, demonstrating high chemical resolution in mixture analysis. The results show that the EMD kernel achieves accurate yield predictions with narrow distributions for both high and low-yield cases, indicating improved performance in distinguishing between different mixture compositions. By incorporating smooth distance metrics, the researchers successfully extended Bayesian optimization techniques from pure components to molecular mixtures, potentially enhancing the efficiency and interpretability of materials property prediction in complex chemical systems.

#### **Project 23: Reliable Surrogate Models of Noisy Data**

This project, awarded first place, asked how reliably BO reaches an optimum when the objective is noisy, and whether encoding a noise assumption into the surrogate helps. The team used the two-dimensional Schwefel function (highly multimodal, with its global minimum geometrically distant from the next-best basins) with additive Gaussian noise, and swept a grid of noise standard deviations (1, 5, 10, 20) and initial design sizes (2, 4, 8, 10) over a 30-evaluation budget with five random seeds. In the BoTorch arm [6], a Gaussian-process surrogate that infers a homoscedastic noise level was compared against the same model with the observation variance pinned to a near-zero constant, scoring each run by the best noise-free function value found. The noise-aware model won in 15 of the 16 grid cells, and its margin grew with noise: negligible at the lowest noise level but 6.6 to 26.5 objective units at the highest. The BayBE arm [15] ran greedy expected-improvement campaigns against a random-recommender baseline over both a restricted range and the wider domain containing the global optimum; there, the team observed that most campaigns settled into local minima far from the global optimum while some high-noise and random runs found the global basin early, so BO’s advantage over random search was not robust at this budget. Their conclusion was that a noise-aware surrogate matters more the noisier the data, but no surrogate rescues a budget too small for the landscape. (No video was submitted; summary prepared from the team’s project page and repository.)

#### **Project 24: ScattBO Benchmark - Bayesian optimisation for materials discovery**

This project presents ScattBO, a Python-based benchmark that simulates a self-driving laboratory (SDL) for

materials discovery. A self-driving laboratory is an autonomous platform that conducts machine learning-augmented experiments to achieve a user-defined objective, such as synthesizing a specific material [52]. The benchmark addresses the challenge that such SDLs can be expensive to run, making intelligent experimental planning essential, while only a few people have access to real SDLs for materials discovery. ScattBO provides an in silico simulation of an SDL where, based on synthesis parameters, the benchmark ‘synthesizes’ a structure, calculates the scattering pattern [53], and compares it to the target structure’s scattering pattern. The benchmark acknowledges that scattering data may not be sufficient to conclusively validate that the target material has been synthesized [54], but can include other types of data as long as they can be simulated. This makes it currently challenging to benchmark Bayesian optimization algorithms for experimental planning tasks in SDLs, and ScattBO fills this gap by providing an accessible simulation environment.

#### **Project 25: Bayesian Optimized De Novo Drug Design for Selective Kinase Targeting**

This project focused on incorporating Bayesian optimization to guide de novo drug design, specifically targeting growth factor receptors for cancer therapeutics. The team built upon the DOCKSTRING paper, Python library, and dataset [55], using a Gaussian process with a Matérn kernel on Morgan fingerprint representations. They employed a graph genetic algorithm to generate SMILES strings guided by the Bayesian optimization output. The researchers explored both selective and promiscuous binding scenarios. For selective binding, they optimized for binding to FGFR1 while penalizing overbinding to other growth factor receptors relative to their median. For promiscuous binding, they maximized the maximum binding affinity across multiple receptors. They found that a sigmoidal penalty function was more effective than simple absolute differences when optimizing against multiple proteins. The team also incorporated a drug-likeness measure (QED) [56] as a penalty in the optimization process, though its effect was limited. Due to time and resource constraints, the project was unable to extensively explore the chemical space or use more accurate binding affinity calculations beyond docking. The authors suggest that future work could incorporate known unknowns through an evasion process, further optimize selective binding, and compare different molecular representations.

#### **Project 26: Multiple-Context Bayesian Optimization**

This project focused on exploring multiple context Bayesian optimization using the BayBE code [15], a Bayesian backend package built on BoTorch with additional features. The primary aim was to investigate transfer learning capabilities by incorporating data from existing campaigns into new optimization tasks. The researchers examined both analytical functions (like the

Hartmann 3D function) and real-world data (direct arylation reaction dataset) to assess the effectiveness of transfer learning in Bayesian optimization. They introduced noise, scaling, shifting, and negation to the analytical functions to simulate realistic scenarios. The results demonstrated that using a small percentage (1-25%) of existing data significantly improved optimization performance compared to the baseline without transfer learning. Interestingly, they found that using larger amounts of data (50-100%) did not necessarily lead to better results and in some cases performed worse than the baseline. The team also explored clustering the existing data and using cluster centroids as source data, which proved effective. For the real-world direct arylation dataset, they analyzed correlations between different reaction temperatures and observed that small amounts of transfer data (1-10%) substantially improved the Bayesian optimization campaign. Overall, the project highlighted the potential of transfer learning in Bayesian optimization while emphasizing that careful selection of the amount and type of transfer data is crucial for optimal performance.

#### **Project 27: How does initial warm-up data influence Bayesian optimization in low-data experimental settings?**

This research project investigated the influence of warm-up sampling methods and dataset sizes on property optimization in low data regimes, specifically focusing on molecular property prediction. The team used the QM9 dataset [57] and selected band gap as the optimization target. They compared two chemically-inspired sampling methods for the warm-up dataset: Morgan fingerprints and MolFormer language model fingerprints. The researchers also referenced the GDB-17 chemical universe database [58] in their background work. The researchers performed dimensionality reduction on the fingerprints using PCA, projecting them into a 2D space for sampling. They conducted experiments to analyze how the warm-up dataset size affects optimization results. The most significant finding was the comparison between Morgan fingerprints and MolFormer fingerprints at a constant data regime of 50 data points. The results showed that MolFormer fingerprints substantially outperformed Morgan fingerprints, suggesting that pre-trained models on large chemical spaces can potentially improve model optimization rates. This study aims to initiate broader discussions on how dataset sizes and sampling methodologies impact final optimization tasks in molecular property prediction.

#### **Project 28: The Impact of Dataset Size on Bayesian Optimization, Insights from the QM9 Dataset**

The Chihuahuas team investigated the critical dataset size threshold required for effective Bayesian Optimization of Graph Convolutional Neural Networks (GCNNs) using the QM9 molecular dataset [40]. The research systematically examined hyperparameter optimization (learning rate, batch size, number of neurons) across dataset sizes ranging from 200 to 2000 molecules to de-

termine minimum viable data requirements for reliable BO performance. Results demonstrated that BO successfully optimized GCNN hyperparameters even with small datasets (200 molecules), though performance improved with larger datasets, achieving lowest losses at 2000 molecules (0.001134 for learning rate optimization). The study revealed a trade-off between dataset size and computational efficiency, with larger datasets requiring more time to reach convergence despite better performance. The research provides a methodological framework for applying BO in data-scarce scenarios common in computational chemistry and materials science, demonstrating that machine learning optimization remains viable even with limited molecular data.

#### **Project 29: A Bayesian Approach to Predict Solubility Parameters**

This project proposed modeling Hansen solubility parameters with Gaussian-process regression on chemical descriptors and sigma profiles, comparing molecular kernels from the GAUCHE framework [59] against an existing Bayesian treatment whose published solubility-parameter compilation was to serve as both benchmark and dataset [60]. The submission stayed at the proposal stage; no dataset was ultimately used and no code or results were published.

#### **Project 30: Active learning for voltammetry waveform design**

This research project focused on using Bayesian Optimization to optimize voltammetry waveforms for sensor performance testing. The researchers employed the scikit-optimize package [61] to implement a human-in-the-loop optimization process, allowing for pauses between optimization iterations to conduct laboratory experiments. All observations were collected in-house: the team measured serotonin concentration prediction error (mean absolute error in nanomolar) in vitro for each candidate waveform using their rapid pulse voltammetry platform [62], and this waveform-performance dataset is recorded directly in the project notebook. The study encoded a four-step pulse waveform into a continuous parameter space for the optimization model. The process began with six randomly generated initial waveforms, which were tested in the lab to obtain performance metrics. These results were used to initialize a Gaussian process model. The optimization then proceeded in batches, with the model suggesting the next best waveforms to test based on previous results. The researchers conducted four total batches, updating the model after each laboratory testing phase. Progress plots were generated to visualize the improvement in sensor error across batches, demonstrating the effectiveness of the approach in finding optimal waveform parameters. The study also explored the interactions between model parameters and the performance metric across multiple dimensions, providing insights into the complex relationship between waveform characteristics and sensor performance.

#### **Project 31: A tutorial on ask/tell mode for Ax**

This project focuses on utilizing the open-source package AX [63] for Bayesian optimization in an asynchronous, human-in-the-loop experimental setting, particularly applied to voltammetry. The researchers developed custom figures and explanations to elucidate the complex interactions within the AX package, emphasizing its user-friendly nature while acknowledging its potential complexity. The study explores various aspects of the optimization process, including strategy selection, surrogate model choice, and search space definition. It demonstrates two primary approaches: a “cold start” scenario, where optimization begins with no prior data, and a “warm start” scenario for situations with existing experimental data. The demonstration data are the team’s own in vitro voltammetry measurements, the same in-house dataset used in project 30 [62]; no external dataset is involved. For the cold start, the researchers showcase the use of sequence generators like Sobol sequences and random generators to initialize the optimization campaign. The project also highlights the package’s ability to handle real-world experimental constraints, such as marking broken electrodes to exclude them from future suggestions. In the warm start scenario, the researchers detail the process of translating existing experimental workflows into the AX framework, emphasizing the package’s automatic tracking of experiments, results, and parameters. The study also notes the inclusion of visualization tools within AX, though these are not described in detail.

#### **Project 32: Efficient Protein Mutagenesis using Bayesian Optimization**

This research project explored the application of Bayesian optimization to protein mutagenesis, specifically aiming to optimize protein binding affinity to fentanyl for potential use in biosensors. The team utilized a pre-trained BERT language model [64], which predicts ligand binding to proteins based on amino acid sequences and SMILES strings. The methodology involved using Bayesian optimization to iteratively select positions and amino acids for mutation, starting with a protein that already had affinity to fentanyl [65]. The researchers focused on residues within 5 angstroms of the ligand as potential mutation sites. The team compared their Bayesian optimization approach to a baseline using random mutations. They analyzed the resulting mutations using PyRosetta’s mutagenesis tool to identify potential new positive interactions and clashes that could alter protein conformation. The structural changes were further investigated using AlphaFold [66].

#### **Project 33: Bayesian Optimization for Hyperspectral Co-heritability Search**

This research project focuses on applying Bayesian Optimization (BO) to the field of crop genetics, specifically for selecting optimal proxy data to train genomic predic-

tion models. The goal is to identify resilient crops that can withstand factors like pests and climate change. The study uses heritability as a measure to determine the best proxy traits for desired traits, treating it as a black-box optimization problem. The researchers compared various BO methods using Gaussian Processes (GPs) against random search for four different desired traits. Surprisingly, they found that all BO methods performed similarly to random search. However, a pre-trained probabilistic model with a mixed multitask kernel outperformed standard BO approaches, identifying better proxy data more quickly. The study concludes that Gaussian Processes may not be the optimal choice for this particular problem, challenging the common use of Matérn kernels and GPs as the default in Bayesian Optimization. The researchers also criticize the widespread use of these methods without proper justification. As part of their contribution, the team has released their dataset [67] for further research in this area.

#### **Project 34: Streamlining Material Discovery - Bayesian Optimization in Thermal Fluid Mixtures**

This project, awarded second place, applied BO to the selection of thermal-fluid mixtures for heat-transfer applications, following the pool-based framing used for nanoporous materials [68]. Each candidate fluid in an in-house library of several hundred mixtures was described by four thermophysical properties (viscosity, density, thermal conductivity and heat capacity), and the objective was the measured heat transfer coefficient. A Gaussian-process surrogate was driven through an active-learning loop that started from ten randomly chosen fluids and ran 50 acquisitions, repeated over 20 replicates and benchmarked against random sampling. BO located the top-ranked fluid within roughly 25 to 30 acquisitions and plateaued near the library maximum of the heat transfer coefficient after about 10 to 20, whereas random sampling was still tens of ranks short of the optimum and several units below the best achievable value after the full budget; the spread across replicates was also markedly narrower for BO, which the team framed as giving experimentalists a more dependable shortlist. Two tuning studies accompanied the comparison: a screen over kernel families and length scales identified the Matérn kernel as the best performer by marginal likelihood, after which a genetic algorithm refined its length scale over 250 generations, and a comparison of acquisition functions found expected improvement slightly ahead of probability of improvement and the upper confidence bound. The underlying property dataset is proprietary and was not released with the code. (No video was submitted; summary prepared from the team’s project page, poster, and repository.)

#### **Project 35: Tutorial for GAUCHE - A Library for Gaussian Processes in Chemistry**

This research project focuses on implementing input warping for Bayesian Optimization within the Gauche library [59], which was previously developed by the team and published at NeurIPS 2023. The primary innovation of Gauche is the introduction of Gaussian process (GP) kernels that enable modeling of discrete entities such as SMILES strings, graphs, and bit vectors, which are common representations in molecular sciences. The motivation behind using Gaussian processes for Bayesian Optimization is their suitability for automated tasks where fine-tuning for each problem is not feasible. GPs offer a good balance between performance and simplicity, with few trainable hyperparameters that can reliably converge on each iteration of the Bayesian Optimization loop. This makes them particularly attractive as surrogate models compared to more complex alternatives like deep neural networks, which might require careful monitoring during training at each iteration. The Gauche library extends the applicability of GPs to discrete input spaces, allowing for Bayesian Optimization over molecular representations. The project team has developed a range of tutorials and applications, including molecular property prediction, protein fitness prediction, and sparse GP regression, all available in the Gauche GitHub repository.

#### **Project 36: Scalable Nonmyopic Bayesian Optimization in Dynamic Cost Settings**

This research project focuses on scalable Bayesian optimization in dynamic settings, addressing limitations of previous approaches that rely on myopic acquisition functions and assume fixed cost structures. The researchers introduce a novel method using non-myopic acquisition functions [69] that incorporate a look-ahead mechanism and dynamic cost functions. The project evaluates the proposed algorithm, named HBE, through two main experimental setups. First, they use synthetic functions across 14 different environments with varying dimensions to test scalability. Second, they apply the method to a real-world protein sequence design problem, aiming to maximize a protein **score** fitness score; in the published version of this work, the scoring oracle is fitted to the ProteinEA fluorescence dataset [70]. The researchers compare their HBE algorithm against six other acquisition functions, including state-of-the-art methods. To enhance practicality, they integrate automatic hyperparameter tuning to reduce the number of optimization parameters. While specific results are not provided in the given context, the approach aims to overcome suboptimal resource allocation in dynamic cost experiments and improve upon existing Bayesian optimization techniques.

### Project 37: The Effects of Post-Modelling Performance Metric Computation on the Efficiency of Bayesian Optimizers

This research project explored modifications to the Bayesian Optimization algorithm to enhance its efficiency in solving real-world problems. The key innovation was changing the order of objective calculation and model fitting. Instead of calculating all objectives first and then modeling each objective, the proposed method only models the essential observations (e.g., yield) and calculates the objectives after the modeling process. This approach was tested ~~using a simulation of a real-world system based on kinetics discovered by Hornel, employing a on a simulated~~ plug flow reactor ~~benchmark of a nucleophilic aromatic substitution reaction from the Summit package [71], which generates yield observations from the kinetic model of Hone et al. [72]; no external experimental dataset was involved.~~ The optimization aimed to maximize ~~both~~ space-time yield and ~~E-factor for minimize the E-factor computed from~~ the predicted yield. The results demonstrated significant time savings by computing objectives after modeling, primarily because only a single model for yield needed to be fitted instead of separate models for space-time yield and E factor. While computational performance improved, there was negligible impact on convergence speed. Interestingly, the research suggests that this new approach might potentially improve the efficiency of selection criteria by allowing the fitting of a prior distribution biased towards regions with more potential based on the objective function’s structure. In summary, the study found that fitting objectives after modeling can lead to substantial increases in computational performance, with the possibility of also improving hypervolume convergence efficiency.

### Project 38: Bayesian methods in symbolic regression

This research project focused on enhancing symbolic regression methods for physical science data sets using Bayesian optimization techniques. The researchers explored two main approaches to improve the performance of symbolic regression algorithms. First, they investigated the impact of assigning varying importance to individual data points under data-limited conditions. Using the Optuna framework, they implemented a tuning process for data point importance, which resulted in improved predictive performance on a benchmark set compared to treating all points equally. This approach could be particularly useful when dealing with experimental data of uncertain quality. Second, they integrated prior knowledge from Wikipedia’s known mathematical equations into the loss function, combining it with the Bayesian Information Criterion to statistically favor established symbolic structures. This method was implemented as a fitness function in the Symbolic Regression Julia library. While the researchers faced challenges in evaluating results due to numerical instability, both approaches showed promise in improving symbolic regression performance by incorporating uncertainty and prior informa-

tion. The team suggests that future work should adhere more closely to the original paper’s methodology, using the Bayesian information prior differences between equations as a probability metric for accepting mutational operations in the symbolic regression process.

### Project 39: Divide and Conquer - Local Gaussian Processes to design Covalent Organic Frameworks for Methane Deliverable Capacity

This research project focuses on improving Bayesian Optimization (BO) for high-dimensional, large-scale datasets, specifically applied to the design of Covalent Organic Frameworks (COFs) for methane storage [68]. The researchers developed a novel approach combining unsupervised clustering with local Gaussian Process (GP) models to enhance BO efficiency in the high data regime. The method begins by using K-means clustering to partition the dataset into distinct clusters. A portion of data from each cluster is then sampled to train separate local GP models. An epsilon-greedy algorithm is employed to determine which GP to train next. The researchers applied this approach to a COF dataset [73] containing over 70,000 2D and 3D structures, assembled in silico from 666 organic linkers and four synthetic routes. The objective was to maximize methane storage performance, measured as deliverable capacity. The results demonstrated that their divide-and-conquer approach with local GP surrogates significantly outperformed a single GP model. While the single GP model reached a maximum deliverable capacity of 2.74 in 60 iterations, the proposed method surpassed this maximum within 5-10 iterations. Both methods retrained the GP model after 20 iterations and used 5% of the data for initial surrogate model training. This research highlights the potential of using local GP surrogates in combination with unsupervised clustering to perform more efficient Bayesian optimization in high-dimensional, large-scale datasets.

### Project 40: Optimizing Chemical Reaction Conditions with Multi-Agent Systems Using LLM and BO

This research project focuses on optimizing Suzuki reaction conditions using Bayesian optimization enhanced by a multi-agent large language model (LLM) system. The target dataset ~~consists of Suzuki reactions, with the goal of optimizing~~, included in the project repository, is drawn from the nanomole-scale Suzuki-Miyaura coupling screen of Perera et al. [74] and tabulates measured yields for 308 combinations of 11 ligands, 7 bases, and 4 solvents; the goal is to optimize reaction ligands, bases, and solvents. The researchers observed that while the target dataset contains unique combinations of conditions, many of the ligands are similar to those in existing reaction data points, suggesting potential for knowledge transfer. The developed system employs a multi-agent LLM approach comprising three domain expert agents and a reaction optimization agent. The process

involves four steps: (1) the reaction optimization agent assigns tags to domain agents, (2) gathers feedback from them, (3) uses this feedback to warm-start initial data point generation and evaluation, and (4) acquires new data points based on existing data. In demo experiments, the multi-agent LLM system demonstrated improved performance, reaching target reaction yields with fewer iterations compared to traditional methods. Additionally, the multi-agent approach showed an ability to avoid local maxima traps, which is particularly valuable given the challenging nature of reaction optimization due to data sparsity. This research highlights the potential of integrating LLM-based multi-agent systems with Bayesian optimization techniques to enhance reaction condition optimization in chemistry.

#### Project 41: Retrieval Augmented Bayesian Optimization

This research project, called RAMBO, combines retrieval-augmented generation with Bayesian optimization to enhance the initial point selection for new optimization tasks. The system leverages existing data from literature or internal databases to identify optimal starting points for Bayesian optimization processes, particularly in the context of chemical reactions. The RAMBO pipeline begins by describing the design space of a reaction to be optimized and determining which reaction parameters should be explored first. It then queries available data to extract the most relevant data points and assembles initial suggestions for the user. The system was demonstrated using a Suzuki-Miura coupling reaction as an example, which involves an extensive combinatorial space of different parameters. RAMBO can describe the design space and limit it to available compounds, providing starting parameters to optimize the reaction. The project includes a demo interface where users can input their reaction of interest, and the React system in the backend queries and extracts relevant data to form a reply containing necessary conditions to initiate the Bayesian optimization process. Users can also explore the reasoning behind the system’s suggestions by analyzing the data and reactions used to generate the final answer.

#### Project 42: Bayesian Optimization and Design of Experiments for Pharmaceutical Formulation (BODoE)

This industry-led project proposed coupling BO-driven design of experiments with an automated laboratory for the formulation of pharmaceutical drug compounds as multi-component mixtures, defining the data schemas and interfaces linking laboratory devices, property-prediction services such as COSMO-SAC [75], and the optimization loop. The submission stayed at the proposal stage: no dataset was used and no code or results were published.

#### Project 43: Bayesian Optimization Awesome List

This awesome list curates a collection of software, tutorials, research papers, and other resources related to Bayesian optimization in materials science and chemistry.

#### Project 44: Rank-based Bayesian Optimization

This research project focused on using ranking models as surrogates in Bayesian optimization for materials discovery, specifically comparing ranking-based and conventional mean squared error (MSE) loss approaches. The key motivation was that in experimental campaigns, finding the best molecule is more important than accurately predicting absolute property values, inspired by work on molecular pool-based active learning [76]. The study employed a pairwise ranking loss (margin ranking loss) and a simple fully-connected multi-layer perceptron with three hidden layers and 100 nodes as the model architecture. Experiments were conducted on multiple datasets, including a solubility dataset (Delaney) and two datasets from Ali et al. with varying roughness [77]. Results showed that the ranking loss model consistently outperformed the MSE-based model and random baseline in Bayesian optimization, acquiring more top-performing candidates within fewer evaluations. Interestingly, all models performed better on the smoother dataset, contrary to expectations. The study also found that model performance did not always correlate with its effectiveness as a surrogate in Bayesian optimization, as evidenced by the MSE model performing worse than the random baseline on the rougher dataset. The research highlights the potential of ranking-based models in overcoming overfitting issues common in Bayesian optimization with limited data points, particularly in materials discovery applications.

#### Project 45: Bayesian Optimization for generality

This research introduces a novel approach to Bayesian Optimization (BO) focused on finding general parameters that perform well across multiple related tasks, a common challenge in fields like chemistry where optimal reaction conditions for various substrates are sought. The researchers formulate this problem within the BO framework, aiming to maximize a general response from a set of multiple related objective functions. The proposed algorithm operates in two steps: first selecting a value in the optimization domain  $X$ , then choosing an individual objective function to optimize for that value. This approach aims to find the most general values in  $X$  without exhaustively evaluating every individual surface, thus minimizing evaluation costs. ~~To~~ All objectives were synthetic rather than drawn from a chemical dataset: to support further research in this area, the team adapted popular test surfaces (Dixon-Price, Branin, and Beale) ~~to be into~~ families of related tasks compatible with their generality problem and made these benchmark problems available on Hugging Face. This contribution allows other researchers to test and develop new algorithms for gener-

Table IV. Classification of the 45 project outputs by primary deliverable. Shares are rounded to whole percentages. Per-project category assignments are listed in the ESI.

| Primary output category   | Count | Share |
|---------------------------|-------|-------|
| Application demonstration | 22    | 49%   |
| Preliminary concept       | 10    | 22%   |
| Benchmark dataset/problem | 6     | 13%   |
| Tutorial/educational      | 4     | 9%    |
| Mature software           | 3     | 7%    |

alized Bayesian Optimization, potentially advancing the field and its applications in natural sciences.

#### IV. Cross-Project Synthesis

The per-project summaries above document *what* each team attempted; here we step back to synthesize patterns across the 45 submissions. Given that the synchronous hackathon spanned two days, many of these projects did not undergo the same rigor and fleshing out as what might be expected of a peer-reviewed manuscript. However, we report qualitative trends and examples.

We classified each submission by its *primary* deliverable into one of five categories: mature software, benchmark dataset/problem, tutorial or educational resource, application demonstration, and preliminary concept (Table IV). The distribution is dominated by application demonstrations (22 of 45), with comparatively few released, reusable software tools (3). The label “mature software” denotes a usable, released tool or a contribution to an established library such as GAUCHE [59].

This deliverable-based grouping cuts across the five special topics under which projects were submitted (Table II): of the 45 project pages, 20 record the open-ended general topic, 9 propose real-world tasks, 5 apply algorithms to benchmark tasks, 4 create tutorials, 2 develop benchmarks, and 5 record no topic. How many submissions landed in the general category surprised the organizers, who would not have known in advance how best to segment the topics [78]; a special topic such as “large language models applied to Bayesian optimization” would likely have been among the largest had it been offered, given the LLM prevalence noted below.

We also assessed code availability and licensing. Of the 45 projects, 42 link to a public code repository; Table V tallies license status per project. Explicit licenses are overwhelmingly permissive (20 MIT, 4 Apache-2.0), 12 projects provide no license file, and 4 link to repositories that were private or removed when checked. Dependency specifications were frequently incomplete, so most outputs are reusable starting points rather than finished tools [79]. The license tabulation script is

Table V. License status of the code repositories linked from Table III, tallied per project so that all 45 projects are accounted for; projects sharing a repository share its license. The CC0-1.0 repository carries its dedication in a form the GitHub API does not auto-classify. “No license file” indicates a public repository lacking a license; “inaccessible” indicates a private or removed repository.

| License                               | Projects  |
|---------------------------------------|-----------|
| MIT                                   | 20        |
| Apache-2.0                            | 4         |
| GPL-3.0                               | 1         |
| CC0-1.0                               | 1         |
| No license file                       | 12        |
| Inaccessible (private/removed)        | 4         |
| No public repository (proposal stage) | 3         |
| <b>Total</b>                          | <b>45</b> |

archived with the manuscript source.

Data provenance was similarly uneven: for roughly a quarter of the submissions, the dataset in use was not evident from the project page or repository alone, so the per-project summaries in the preceding section state each project’s data source explicitly wherever it is not obvious. The patterns are few. Some projects retrospectively queried published experimental datasets, for example the direct arylation yields of Shields et al. [30] in project 15 and the Suzuki-Miyaura screen of Perera et al. [74] in project 40. Others generated objective values on demand from physics-based simulators or synthetic test functions (projects 6, 36, 37, and 45), collected new measurements in-house (projects 30 and 31), or, for proposal-stage submissions, used no data at all (projects 19, 29, and 42). Even where a published dataset was used, it was more often fetched at run time or referenced from within the code than deposited in the repository: project 15, for instance, pulls its benchmark directly from the EDBO+ repository [31]. Archiving a copy of the queried dataset alongside the project code, with its provenance stated in the README, would make such outputs considerably easier to audit and reuse.

Methodologically, a common template emerged: frame a chemistry or materials task as sequential optimization of an expensive objective with a Gaussian-process surrogate and an acquisition function, most often built on BoTORCH/Ax, sometimes through domain wrappers such as BAYBE [80] or GAUCHE [59]. Teams differed most in input representation, in whether they compared against simple baselines such as random search, and in whether they targeted real experimental data or retrospective benchmarks. Participants ranged from students to senior researchers across academia, industry, and government, and mixed-expertise teams pairing domain scientists with BO practitioners appeared to be an asset [81, 82]. Roughly one in six projects (6 of 45 under a strict definition, 8 under a broader one)

incorporated large language models, mirroring dedicated LLM-for-chemistry hackathons [2, 83] and motivating the outlook in Section VI.

## V. Lessons Learned

### A. Organizational lessons

The reflections below draw in part on a recorded interview about how this event was designed and run, given by the lead organizer to the organizers of a subsequent machine-learning-for-microscopy hackathon as they prepared their own event [78, 84].

Several organizational choices worked well and would be retained. The scope was kept deliberately open-ended, following hackathon guides that warn against over-constraining the problem: Bayesian optimization for chemistry and materials, loosely segmented into five special topics, was broad enough to admit creative proposals while remaining coherent, and the open-ended general topic indeed drew the most submissions (Section IV) [78]. Extensive pre-event preparation (orientation materials, an introduction to BO, curated tooling, and a Python refresher delivered through GitHub Classroom with automated feedback) helped level a wide range of backgrounds (Figure 3) and set a shared baseline of what teammates could expect of one another; the real-world-task topic complemented this by giving domain experts without a programming background a way to contribute meaningfully [78]. A persistent virtual venue (Gather Town) sustained informal interaction and mentor access in a fully online setting, and a transparent pairwise-comparison judging system (Gavel) scaled expert evaluation across many projects with modest per-judge effort (35 judges cast 319 votes). Archiving every output to GitHub and consolidating them in a single Zenodo record produced a durable, citable community resource, and several outputs have persisted past the event, including contributions folded into the maintained GAUCHE library and a community-curated resource list (project 43) [83]. These choices align with best practices identified across scientific hackathons: strong onboarding, structured mentoring, and disciplined post-event archiving [81, 82, 85]. Consulting the organizers of the preceding LLM hackathon directly, beyond reading their retrospective, also helped anticipate problems [2, 78].

Other aspects proved more difficult. Some teams that registered before the event did not converge into active projects, and six teams recorded no closing video, which made their outputs harder to capture; pre-event team formation and mentor engagement are reportedly among the most fragile parts of virtual hackathons [86]. Our experience was consistent with this: the project list did not populate itself. Many prospective leads wanted to discuss an idea privately with an organizer

before posting it, and the first ten to fifteen proposals were seeded through persistent one-on-one outreach; once that critical mass made the format visible, later teams followed the pattern and filled out the final roster of 45 [78]. The event was announced only about six weeks in advance, and much of its value came from participants blocking out the same two days, which concentrates attention the way travel to a conference does [78]. Communication infrastructure was also harder than expected: single-channel Slack guest invites, chosen so that discussions would remain searchable in the organizers’ workspace, proved cumbersome to administer across workspaces [78]. The missing videos did not track project quality: the first- and second-place projects were among the six without one. The compressed two-day window limited validation depth, and the reproducibility of participant code varied widely (Section IV). In hindsight, we would (i) announce the event and begin team formation earlier, with structured virtual meet-and-greets one to two weeks in advance; (ii) provide explicit reproducibility and licensing checklists (pinned dependencies, a runnable notebook, and an appropriate license) to raise the floor on code quality; (iii) request a short recorded or written summary from every team so that all outputs are captured uniformly; (iv) communicate evaluation expectations to participants more explicitly; (v) host discussion in a dedicated Slack or Discord channel for the months around the event rather than through guest invites; and (vi) run a follow-up survey of participants, since without one we cannot report whether the event changed how participants use BO in their own research. Judging used holistic pairwise comparison rather than a numeric rubric (see Hackathon Details and Setup for the judging question); the submission and authorship requirements that framed evaluation are documented on the hackathon website and archived in the Zenodo record.

### B. Scientific lessons: strengths and limitations of Bayesian optimization

Taken together, the projects illustrate both the appeal and the practical pitfalls of BO in chemistry and materials, in line with the broader literature. On the strengths side, teams repeatedly exploited BO’s sample efficiency in low-data regimes, recovering strong candidates within a handful of iterations relative to design-of-experiments or random baselines [87, 88]. The submissions collectively demonstrated BO’s flexibility, handling continuous, categorical, mixed, molecular, multi-objective, multi-fidelity, contextual, batched, and cost-aware problems, including successful prototypes of multi-objective search and preference-based, human-in-the-loop workflows [89]. Specialized kernels and representations (e.g., for molecules) further extended BO to structured inputs [59].

The projects also surfaced well-known limitations. Default BO was not always superior to simple baselines: two teams reported little advantage

over random search on their tasks (projects 33 and 1), and while these two-day comparisons were not run at benchmark rigor, they underscore the value of always including a simple baseline [90]. Results were sensitive to molecular representation, featurization, and kernel choice; warm-start data quality, observation noise, and high-dimensional search spaces were recurring pain points, for which noise-robust and heteroscedastic-aware methods are relevant remedies [91, 92]. Acquisition-function optimization itself mattered: one of the top-ranked projects (project 3) showed that investing more computation in the acquisition optimizer changed observed performance, connecting to ongoing work on scalable and high-dimensional acquisition strategies [46, 93]. Finally, retrospective benchmark gains do not guarantee wet-laboratory success, and reproducing comparisons across studies remains difficult [90]. Where teams integrated LLMs, the models lowered access barriers and helped specify spaces or retrieve priors, but they added another uncertain component that can degrade on complex objectives and requires validation against non-LLM baselines [94].

On the recurring question of which framework works best for a given noise regime, these projects support only a narrow observation: in project 23, modeling the observation noise in the surrogate became more valuable as the noise grew. More broadly, deployment-oriented libraries such as BAYBE and BOFIRE target the experimental setting specifically [80, 95]; a general ranking of frameworks by noise regime is beyond what a two-day event can establish.

## VI. Future Opportunities

Building on the hackathon outputs and the trends above, we highlight several promising directions for BO in chemistry and materials discovery. *Autonomous and self-driving laboratories* are a natural home for BO, closing the loop between algorithmic suggestion and automated experimentation [96, 97]. *Multi-fidelity optimization* that fuses cheap simulations or proxy measurements with expensive experiments can substantially reduce cost, provided its assumptions are handled carefully [98, 99]. *Uncertainty-aware experimental planning* and active learning under design constraints can make campaigns more sample-efficient and trustworthy [100].

The large share of hackathon projects that already incorporated language models points to *foundation-model- and LLM-assisted BO* as a particularly active frontier for design-space specification, prior construction, and warm-starting; reported evidence on whether LLM-based surrogates and priors outperform strong classical baselines is mixed, and benchmarking such hybrids fairly (separating genuine generalization from memorized

results) is itself an open problem [94, 101]. Other opportunities include *preference-based and human-in-the-loop BO* for subjective or hard-to-quantify objectives [102, 103]; *multi-objective materials design* with scalable hypervolume-based acquisitions [93]; *robust BO under noisy experimental data* [91, 92]; *community benchmark development* to enable reproducible, apples-to-apples comparisons [71, 90]; and *domain-specific, deployment-ready BO tools* for synthesis and processing optimization that meet experimentalists where they work, including a growing set of graphical and web-based BO interfaces that remove coding barriers [31, 80, 95, 104, 105]. Collectively, these directions echo the broader perspective that Bayesian methods are becoming a core component of the autonomous materials-discovery toolkit [106].

## Data Availability

The code and materials associated with the projects in this hackathon are archived on Zenodo at [10.5281/zenodo.20496810](https://zenodo.org/record/105281/versions/10).

## Conflicts of interest

There are no conflicts of interest to declare.

## Acknowledgements

We express a huge thanks to Kevin Jablonka and Ben Blaiszik (latter a co-author of this paper) and others did phenomenal work paving the way with the LLMs for Chemistry and Materials Hackathon (2023). Much of the setup and many of the patterns of the Bayesian optimization hackathon was either directly replicated based off of or inspired by the LLM hackathon.

This work was undertaken thanks in part to funding provided to the University of Toronto’s Acceleration Consortium from the Canada First Research Excellence Fund (CFREF-2022-00042). We kindly acknowledge Tiago L. P. Galvão for providing Al alloy data from the CORDATA database. Can Özkan acknowledges the VIPCOAT project (Virtual Open Innovation Platform for Active Protective Coatings Guided by Modelling and Optimisation) funded by the Horizon 2020 research and innovation programme of the European Union by grant agreement no. 952903. Ankur Kumar Gupta acknowledges support by the U.S. Department of Energy, Office of Science, Basic Energy Sciences, Materials Sciences and Engineering Division under Contract No. DE-AC02-05CH11231, FWP No. DAC-LBL-Long. Andy S. Anker acknowledges support by Novo Nordisk Foundation grant NNF23OC0081359. Alexander Wiczorek acknowledges funding from the Strategic Focus Area –Advanced Manufacturing ([SFA-AMSFA-AM](#)) through the project Ad-

vancing manufacturability of hybrid ~~organic-inorganic~~ organic-inorganic semiconductors for large area optoelectronics (AMYS).

Finally, we express thanks to all participants in the hackathon. The hackathon was only possible and mean-

ingful due to the time dedicated by many in the projects themselves.

## References

- [1] S. G. Baird and T. D. Sparks, Building a “Hello World” for self-driving labs: The Closed-loop Spectroscopy Lab Light-mixing demo, *STAR Protocols* **4**, 102329 (2023).
- [2] K. M. Jablonka, Q. Ai, A. Al-Feghali, S. Badhwar, J. D. Bocarsly, A. M. Bran, S. Bringuier, L. C. Brinson, K. Choudhary, D. Circi, S. Cox, W. A. De Jong, M. L. Evans, N. Gastellu, J. Genzling, M. V. Gil, A. K. Gupta, Z. Hong, A. Imran, S. Kruschwitz, A. Labarre, J. Lála, T. Liu, S. Ma, S. Majumdar, G. W. Merz, N. Moitessier, E. Moubarak, B. Mouriño, B. Pelkie, M. Pieler, M. C. Ramos, B. Ranković, S. G. Rodrigues, J. N. Sanders, P. Schwaller, M. Schwarting, J. Shi, B. Smit, B. E. Smith, J. Van Herck, C. Völker, L. Ward, S. Warren, B. Weiser, S. Zhang, X. Zhang, G. A. Zia, A. Scourtas, K. J. Schmidt, I. Foster, A. D. White, and B. Blaiszik, 14 examples of how LLMs can transform materials science and chemistry: a reflection on a large language model hackathon, *Digital Discovery* **2**, 1233 (2023).
- [3] G. Mulholland and B. Meredig, Hackathon aims to solve materials problems, *MRS Bulletin* **40**, 366 (2015).
- [4] T. D. Sparks, F. E. Curtis, D. C. Fredrickson, and N. A. Benedek, Insights and Innovations from the SSMCDAT 2023: Bridging Solid-State Materials Chemistry and Data Science, *Chemistry of Materials* **36**, 5293 (2024), publisher: American Chemical Society.
- [5] K. Kandasamy, K. R. Vysyaraju, W. Neiswanger, B. Paria, C. R. Collins, J. Schneider, B. Poczos, and E. P. Xing, Tuning hyperparameters without grad students: Scalable and robust bayesian optimisation with dragonfly, *Journal of Machine Learning Research* **21**, 1 (2020).
- [6] M. Balandat, B. Karrer, D. R. Jiang, S. Daulton, B. Letham, A. G. Wilson, and E. Bakshy, BoTorch: A Framework for Efficient Monte-Carlo Bayesian Optimization, in *Advances in Neural Information Processing Systems* **33** (2020).
- [7] S. Daulton, M. Balandat, and E. Bakshy, Parallel bayesian optimization of multiple noisy objectives with expected hypervolume improvement, *Advances in Neural Information Processing Systems* **34**, 2187 (2021).
- [8] M. Poloczek, J. Wang, and P. Frazier, Multi-information source optimization, *Advances in neural information processing systems* **30** (2017).
- [9] B. Shahriari, K. Swersky, Z. Wang, R. P. Adams, and N. De Freitas, Taking the human out of the loop: A review of bayesian optimization, *Proceedings of the IEEE* **104**, 148 (2015).
- [10] J. Snoek, H. Larochelle, and R. P. Adams, Practical bayesian optimization of machine learning algorithms, *Advances in neural information processing systems* **25** (2012).
- [11] E. S. Muckley, J. E. Saal, B. Meredig, C. S. Roper, and J. H. Martin, Interpretable models for extrapolation in scientific machine learning, *Digital Discovery* **2**, 1425 (2023).
- [12] J. Bergstra and Y. Bengio, Random search for hyperparameter optimization, *The journal of machine learning research* **13**, 281 (2012).
- [13] M. Li, M. J. McCourt, A. J. Galante, and P. W. Leu, Bayesian optimization of nanophotonic electromagnetic shielding with very high visible transparency, *Optics Express* **30**, 33182 (2022).
- [14] J. Kim and S. Choi, Bayeso: A bayesian optimization framework in python, *Journal of Open Source Software* **8**, 5320 (2023).
- [15] M. Fitzner, A. Šošić, A. Hopp, and A. Lee, BayBE – a bayesian back end for design of experiments, <https://github.com/emdgroup/baybe> (2022), accessed: 2024-02-22.
- [16] T. L. Galvão, I. Ferreira, A. Kuznetsova, G. Novell-Leruth, C. Song, C. Feiler, S. V. Lamaka, C. Rocha, F. Maia, M. L. Zheludkevich, *et al.*, Cordata: An open data management web application to select corrosion inhibitors, *npj Materials Degradation* **6**, 48 (2022).
- [17] H. Moriwaki, Y.-S. Tian, N. Kawashita, and T. Takagi, Mordred: a molecular descriptor calculator, *Journal of cheminformatics* **10**, 4 (2018).
- [18] J. L. Durant, B. A. Leland, D. R. Henry, and J. G. Nourse, Reoptimization of mdl keys for use in drug discovery, *Journal of chemical information and computer sciences* **42**, 1273 (2002).
- [19] G. Landrum, Rdkit documentation, Release **1**, 4 (2013).
- [20] Z. Zhou, S. Kearnes, L. Li, R. N. Zare, and P. Riley, Optimization of molecules via deep reinforcement learning, *Scientific reports* **9**, 10752 (2019).
- [21] H. Furukawa, K. E. Cordova, M. O’Keeffe, and O. M. Yaghi, The chemistry and applications of metal-organic frameworks, *Science* **341**, 1230444 (2013).
- [22] D. Y. Heo, H. H. Do, S. H. Ahn, and S. Y. Kim, Metal-organic framework materials for perovskite solar cells, *Polymers* **12**, 2061 (2020).
- [23] D. DeSantis, J. A. Mason, B. D. James, C. Houchins, J. R. Long, and M. Veenstra, Techno-economic analysis of metal-organic frameworks for hydrogen and natural gas storage, *Energy & Fuels* **31**, 2024 (2017).
- [24] F. L. Oliveira, C. Cleeton, R. Neumann Barros Ferreira, B. Luan, A. H. Farmahini, L. Sarkisov, and M. Steiner, Crafted: An exploratory database of simulated adsorption isotherms of metal-organic frameworks, *Scientific Data* **10**, 230 (2023).
- [25] M. Dusselier and M. E. Davis, Small-pore zeolites: synthesis and catalysis, *Chemical reviews* **118**, 5265 (2018).
- [26] A. J. Mallette, K. Shilpa, and J. D. Rimer, The current understanding of mechanistic pathways in zeolite crystallization, *Chemical Reviews* **124**, 3416 (2024).
- [27] M. Altamirano, F.-X. Briol, and J. Knoblauch, Ro-

- bust and conjugate gaussian process regression, arXiv preprint arXiv:2311.00463 (2023).
- [28] S. Ament, A. Witte, N. Garg, and J. Kusuma, Sustainable concrete via bayesian optimization, 2023, in URL <https://arxiv.org/abs/2310.18288>. *NeurIPS 2023 Workshop on Adaptive Experimentation in the Real World*.
  - [29] M. Ansari and A. D. White, Learning Peptide Properties with Positive Examples Only, *bioRxiv* (2023).
  - [30] B. J. Shields, J. Stevens, J. Li, M. Parasram, F. Damani, J. I. M. Alvarado, J. M. Janey, R. P. Adams, and A. G. Doyle, Bayesian reaction optimization as a tool for chemical synthesis, *Nature* **590**, 89 (2021).
  - [31] J. A. Garrido Torres, S. H. Lau, P. Anchuri, J. M. Stevens, J. E. Tabora, J. Li, A. Borovika, R. P. Adams, and A. G. Doyle, A multi-objective active learning platform and web app for reaction optimization, *Journal of the American Chemical Society* **144**, 19999 (2022).
  - [32] H. Mahmoudi, M. Mahmoudi, O. Doustdar, H. Jahangiri, A. Tsolakis, S. Gu, and M. LechWyszynski, A review of fischer tropesch synthesis process, mechanism, surface chemistry and catalyst formulation, *Biofuels Engineering* **2**, 11 (2017).
  - [33] G. Lozano-Blanco, J. W. Thybaut, K. Surla, P. Galtier, and G. B. Marin, Single-event microkinetic model for fischer-tropsch synthesis on iron-based catalysts, *Industrial & engineering chemistry research* **47**, 5879 (2008).
  - [34] A. Chakkingal, P. Janssens, J. Poissonnier, M. Virginie, A. Y. Khodakov, and J. W. Thybaut, Multi-output machine learning models for kinetic data evaluation: A fischer-tropsch synthesis case study, *Chemical Engineering Journal* **446**, 137186 (2022).
  - [35] Z. J. Lin, R. Astudillo, P. Frazier, and E. Bakshy, Preference exploration for efficient bayesian optimization with multiple outcomes, in *International Conference on Artificial Intelligence and Statistics* (PMLR, 2022) pp. 4235–4258.
  - [36] H. Bellamy, A. A. Rehim, O. I. Orhobor, and R. King, Batched bayesian optimization for drug design in noisy environments, *Journal of Chemical Information and Modeling* **62**, 3970 (2022).
  - [37] K. Wu and G.-W. Wei, Quantitative toxicity prediction using topology based multitask deep neural networks, *Journal of chemical information and modeling* **58**, 520 (2018).
  - [38] D. Chen, K. Gao, D. D. Nguyen, X. Chen, Y. Jiang, G.-W. Wei, and F. Pan, Algebraic graph-assisted bidirectional transformers for molecular property prediction, *Nature communications* **12**, 3521 (2021).
  - [39] D. Rogers and M. Hahn, Extended-connectivity fingerprints, *Journal of chemical information and modeling* **50**, 742 (2010).
  - [40] R. Ramakrishnan, P. O. Dral, M. Rupp, and O. A. von Lilienfeld, Quantum chemistry structures and properties of 134 kilo molecules, *Scientific Data* **1**, 140022 (2014).
  - [41] K. Yang, M. Emmerich, A. Deutz, and T. Bäck, Multi-objective bayesian global optimization using expected hypervolume improvement gradient, Swarm and evolutionary computation **44**, 945 (2019).
  - [42] J. Knowles, Parego: A hybrid algorithm with on-line landscape approximation for expensive multiobjective optimization problems, *IEEE transactions on evolutionary computation* **10**, 50 (2006).
  - [43] F. Rapp and M. Roth, Quantum Gaussian process regression for Bayesian optimization, *Quantum Machine Intelligence* **6**, 5 (2024).
  - [44] J. Stetefeld, S. A. McKenna, and T. R. Patel, Dynamic light scattering: a practical guide and applications in biomedical sciences, *Biophysical reviews* **8**, 409 (2016).
  - [45] F. Sorourifar, T. Banker, and J. A. Paulson, Accelerating Black-Box Molecular Property Optimization by Adaptively Learning Sparse Subspaces, *arXiv preprint arXiv:2401.01398* (2024).
  - [46] N. Maus, H. T. Jones, J. S. Moore, M. J. Kusner, J. Bradshaw, and J. R. Gardner, Local Latent Space Bayesian Optimization over Structured Inputs, *arXiv preprint arXiv:2201.11872* (2023).
  - [47] Z. Wu, B. Ramsundar, E. N. Feinberg, J. Gomes, C. Geniesse, A. S. Pappu, K. Leswing, and V. Pande, MoleculeNet: a benchmark for molecular machine learning, *Chemical Science* **9**, 513 (2018).
  - [48] M. J. Kusner, B. Paige, and J. M. Hernández-Lobato, Grammar variational autoencoder, in *Proceedings of the 34th International Conference on Machine Learning*, Vol. 70 (2017) pp. 1945–1954.
  - [49] J. Hachmann, R. Olivares-Amaya, S. Atahan-Evrenk, C. Amador-Bedolla, R. S. Sánchez-Carrera, A. Gold-Parker, L. Vogt, A. M. Brockway, and A. Aspuru-Guzik, The Harvard Clean Energy Project: large-scale computational screening and design of organic photovoltaics on the world community grid, *The Journal of Physical Chemistry Letters* **2**, 2241 (2011).
  - [50] H. B. Moss and R.-R. Griffiths, Gaussian process molecule property prediction with flowmo, *arXiv preprint arXiv:2010.01118* (2020).
  - [51] C. J. Hargreaves, M. S. Dyer, M. W. Gaultois, V. A. Kurlin, and M. J. Rosseinsky, The earth mover’s distance as a metric for the space of inorganic compositions, *Chemistry of Materials* **32**, 10610 (2020).
  - [52] N. J. Szymanski, Y. Zeng, H. Huo, C. J. Bartel, H. Kim, A. Parija, N. Shoghi, S. Rigas, J. B. Neaton, G. Ceder, K. A. Persson, F. M. Toma, and J. M. Gregoire, An autonomous laboratory for the accelerated synthesis of novel materials, *Nature* **624**, 86 (2023).
  - [53] F. L. Johansen, A. S. Anker, T. L. Christiansen, L. S. Voss, P. Juhás, S. J. Billinge, and K. M. Jensen, A GPU-Accelerated Open-Source Python Package for Calculating Powder Diffraction, Small-Angle-, and Total Scattering with the Debye Scattering Equation, *Journal of Open Source Software* **9**, 6024 (2024).
  - [54] J. Leeman, A. A. Kunitsa, K. S. Vecchio, G. Ceder, and A. G. Kusne, Challenges in High-Throughput Inorganic Materials Prediction and Autonomous Synthesis, *PRX Energy* **3**, 011002 (2024).
  - [55] M. García-Ortegón, G. N. C. Simm, A. J. Tripp, J. M. Hernández-Lobato, A. Bender, and S. Bacallado, Dockstring: Easy molecular docking yields better benchmarks for ligand design, *Journal of Chemical Information and Modeling* **62**, 3486 (2022).
  - [56] G. R. Bickerton, G. V. Paolini, J. Besnard, S. Muresan, and A. L. Hopkins, Quantifying the Chemical Beauty of Drugs, *Nature Chemistry* **4**, 90 (2012).
  - [57] R. Ramakrishnan, P. O. Dral, M. Rupp, and O. A. von Lilienfeld, Quantum chemistry structures and properties of 134 kilo molecules, *Scientific Data* **1**, 140022 (2014).
  - [58] L. Ruddigkeit, R. van Deursen, L. C. Blum, and J.-

- L. Reymond, Enumeration of 166 billion organic small molecules in the chemical universe database gdb-17, *J. Chem. Inf. Model.* **52**, 2864 (2012).
- [59] R.-R. Griffiths, L. Klärner, H. Moss, A. Ravuri, S. Truong, Y. Du, S. Stanton, *et al.*, Gauche: A library for gaussian processes in chemistry, in *Advances in Neural Information Processing Systems*, Vol. 36 (2024).
- [60] B. Sanchez-Lengeling, L. M. Roch, J. G. Perea, S. Langner, C. Brabec, and A. Aspuru-Guzik, A Bayesian Approach to Predict Solubility Parameters, *Advanced Theory and Simulations* **2**, 1800069 (2018).
- [61] T. Head, MechCoder, G. Louppe, I. Shcherbatyi, fcharas, D. Cournapeau, F. Pedregosa, and A. Mueller, *scikit-optimize/scikit-optimize: v0.9.0* (2021).
- [62] C. S. Movassaghi, K. A. Perrotta, H. Yang, R. Iyer, X. Cheng, M. Dagher, M. A. Fillol, and A. M. Andrews, Simultaneous serotonin and dopamine monitoring across timescales by rapid pulse voltammetry with partial least squares regression, *Analytical and Bioanalytical Chemistry* **413**, 6747 (2021).
- [63] E. Bakshy, Ax: A platform for adaptive experimentation, GitHub repository (2018).
- [64] A. E. Blanchard, J. Gounley, D. Bhowmik, M. Chandra Shekar, I. Lyngaas, S. Gao, J. Yin, A. Tsaris, F. Wang, and J. Glaser, Language models for the prediction of sars-cov-2 inhibitors, *The International Journal of High Performance Computing Applications* **36**, 587 (2022).
- [65] M. J. Bick, P. J. Greisen, K. J. Morey, M. S. Antunes, D. La, B. Sankaran, L. Reymond, K. Johnsson, J. I. Medford, and D. Baker, Computational design of environmental sensors for the potent opioid fentanyl, *Elife* **6**, e28909 (2017).
- [66] J. Jumper, R. Evans, A. Pritzel, T. Green, M. Figurnov, O. Ronneberger, K. Tunyasuvunakool, R. Bates, A. Žídek, A. Potapenko, *et al.*, Highly accurate protein structure prediction with alphafold, *nature* **596**, 583 (2021).
- [67] STAIR Lab, coh2: Conversational hf interaction dataset, <https://huggingface.co/datasets/stair-lab/coh2> (2025), accessed: 2025-06-15.
- [68] A. Deshwal, C. M. Simon, and J. R. Doppa, Bayesian Optimization of Nanoporous Materials, *Molecular Systems Design & Engineering* **6**, 1066 (2021).
- [69] S. Jiang, D. Jiang, M. Balandat, B. Karrer, J. Gardner, and R. Garnett, Efficient nonmyopic bayesian optimization via one-shot multi-step trees, in *Advances in Neural Information Processing Systems*, Vol. 33 (2020) pp. 18039–18049.
- [70] S. T. Truong, D. Q. Nguyen, W. Neiswanger, R.-R. Griffiths, S. Ermon, N. Haber, and S. Koyejo, Neural non-myopic Bayesian optimization in dynamic cost settings, *arXiv preprint arXiv:2601.06505* (2026).
- [71] K. C. Felton, J. G. Rittig, and A. A. Lapkin, Summit: benchmarking machine learning methods for reaction optimisation, *Chemistry-Methods* **1**, 116 (2021).
- [72] C. A. Hone, N. Holmes, G. R. Akién, R. A. Bourne, and F. L. Muller, Rapid multistep kinetic model generation from transient flow data, *Reaction Chemistry & Engineering* **2**, 103 (2017).
- [73] R. Mercado, R.-S. Fu, A. V. Yakutovich, L. Talirz, M. Haranczyk, and B. Smit, In silico design of 2d and 3d covalent organic frameworks for methane storage applications, *Chemistry of Materials* **30**, 5069 (2018).
- [74] D. Perera, J. W. Tucker, S. Brahmabhatt, C. J. Helal, A. Chong, W. Farrell, P. Richardson, and N. W. Sach, A platform for automated nanomole-scale reaction screening and micromole-scale synthesis in flow, *Science* **359**, 429 (2018).
- [75] MQS, Cosmo-sac: A powerful tool for solvent selection in pharmaceutical drug development, [https://blog.mqs.dk/posts/10\\_cosmo/10\\_cosmo/](https://blog.mqs.dk/posts/10_cosmo/10_cosmo/) (2024), accessed: 2024-05-28.
- [76] D. E. Graff, E. I. Shakhnovich, and C. W. Coley, Accelerating high-throughput virtual screening through molecular pool-based active learning, *Chemical Science* **12**, 7866 (2021).
- [77] M. Aldeghi, F. Häse, R. J. Hickman, I. Tamblyn, and A. Aspuru-Guzik, Roughness of Molecular Property Landscapes and Its Impact on Modellability, *Journal of Chemical Information and Modeling* **62**, 4660 (2022).
- [78] S. G. Baird, Y. Liu, and U. Pratiush, Interview 2: Introduction to hackathon, <https://kaliningroup.github.io/mic-hackathon/videos-slides/> (2024), video interview recorded in preparation for the 2024 hackathon on machine learning for electron and scanning probe microscopy, University of Tennessee, Knoxville. Accessed: 2026-09-23.
- [79] A. Nolte, I.-A. Chounta, and J. D. Herbsleb, What Happens to All These Hackathon Projects? Identifying Factors To Promote Hackathon Project Continuation, *Proceedings of the ACM on Human-Computer Interaction* **4**, 1 (2020).
- [80] M. Fitzner, A. Šošić, A. Hopp, and A. Lee, BayBE – a bayesian back end for design of experiments, <https://github.com/emdgroupp/baybe> (2022), accessed: 2024-05-28.
- [81] M. T. Valentine and R. M. Robertson-Anderson, Hackative Matter: Data-driven Discovery through Hackathon-based Cross-disciplinary Coding, *Soft Matter* **21**, 5381 (2025).
- [82] B. Busby, A. Dillman, C. L. Simpson, *et al.*, Building Genomic Analysis Pipelines in a Hackathon Setting with Bioinformatician Teams: DNA-seq, Epigenomics, Metagenomics and RNA-seq, *bioRxiv* **10.1101/018085** (2015).
- [83] Y. Zimmermann, A. Bazgir, Z. Afzal, *et al.*, Reflections from the 2024 Large Language Model (LLM) Hackathon for Applications in Materials Science and Chemistry, *arXiv* **10.48550/arXiv.2411.15221** (2024).
- [84] U. Pratiush, A. Houston, K. Barakati, *et al.*, Mic-hackathon 2024: hackathon on machine learning for electron and scanning probe microscopy, *Machine Learning: Science and Technology* **6**, 040701 (2025).
- [85] A. Stoltzfus, M. Rosenberg, H. Lapp, A. Budd, K. Cranston, E. Pontelli, S. Oliver, and R. A. Vos, Community and Code: Nine Lessons from Nine NESCent Hackathons, *F1000Research* **6**, 786 (2017).
- [86] M. G. Diemar, C. A. M. Krul, M. Teunis, F. Busquet, H. Kandarov, J. D. Zajac, M. Vinken, and E. L. Roggen, Report of the First ONTOX Hackathon: The Use of Artificial Intelligence in Toxicology, *Alternatives to Laboratory Animals* **53**, 42 (2025).
- [87] P. I. Frazier, A Tutorial on Bayesian Optimization, *arXiv* (2018).
- [88] B. Shahriari, K. Swersky, Z. Wang, R. P. Adams, and N. de Freitas, Taking the Human out of the Loop: A Review of Bayesian Optimization, *Proceedings of the*

- IEEE **104**, 148 (2016).
- [89] Z. J. Lin, R. Astudillo, P. Frazier, and E. Bakshy, Preference Exploration for Efficient Bayesian Optimization with Multiple Outcomes, in *Proceedings of The 25th International Conference on Artificial Intelligence and Statistics* (PMLR, 2022) pp. 4235–4258.
  - [90] Q. Liang, A. E. Gongora, Z. Ren, A. Tihihonen, Z. Liu, S. Sun, J. R. Deneault, D. Bash, F. Mekki-Berrada, S. A. Khan, K. Hippalgaonkar, B. Maruyama, K. A. Brown, J. Fisher III, and T. Buonassisi, Benchmarking the Performance of Bayesian Optimization across Multiple Experimental Materials Science Domains, *npj Computational Materials* **7**, 188 (2021).
  - [91] M. Aldeghi, F. Häse, R. J. Hickman, I. Tamblyn, and A. Aspuru-Guzik, Golem: an Algorithm for Robust Experiment and Process Optimization, *Chemical Science* **12**, 14792 (2021).
  - [92] M. Altamirano, F.-X. Briol, and J. Knoblauch, Robust and conjugate Gaussian process regression, *arXiv preprint arXiv:2311.00463* (2023).
  - [93] S. Ament, A. Witte, N. Garg, and J. Kusuma, Sustainable concrete via bayesian optimization, in *NeurIPS 2023 Workshop on Adaptive Experimentation in the Real World* (2023).
  - [94] A. Kristiadi, F. Strieth-Kalthoff, M. Skreta, P. Poupart, A. Aspuru-Guzik, and G. Pleiss, A Sober Look at LLMs for Material Discovery: Are They Actually Good for Bayesian Optimization over Molecules?, *arXiv* [10.48550/arXiv.2402.05015](https://arxiv.org/abs/10.48550/arXiv.2402.05015) (2024).
  - [95] J. P. Dürholt, T. S. Asche, J. Kleinekorte, *et al.*, BoFire: Bayesian Optimization Framework Intended for Real Experiments, *arXiv* [10.48550/arXiv.2408.05040](https://arxiv.org/abs/10.48550/arXiv.2408.05040) (2024).
  - [96] G. Tom, S. P. Schmid, S. G. Baird, Y. Cao, K. Darvish, H. Hao, S. Lo, S. Pablo-García, E. M. Rajaonson, M. Skreta, N. Yoshikawa, S. Corapi, G. D. Akkoc, F. Strieth-Kalthoff, M. Seifrid, and A. Aspuru-Guzik, Self-Driving Laboratories for Chemistry and Materials Science, *Chemical Reviews* **124**, 9633 (2024).
  - [97] M. Abolhasani and E. Kumacheva, The Rise of Self-driving Labs in Chemical and Materials Sciences, *Nature Synthesis* **2**, 483 (2023).
  - [98] V. Sabanza-Gil, A. Bailly, *et al.*, Best Practices for Multi-fidelity Bayesian Optimization in Materials and Molecular Research, *Nature Computational Science* **5**, 572 (2025).
  - [99] E. Judge, T. Ren, *et al.*, Applying Multi-Fidelity Bayesian Optimization in Chemistry: Open Challenges and Major Considerations, *arXiv* [10.48550/arXiv.2409.07190](https://arxiv.org/abs/10.48550/arXiv.2409.07190) (2024).
  - [100] A. K. Y. Low, F. Mekki-Berrada, A. Gupta, A. Ostudin, J. Xie, E. Vissol-Gaudin, Y.-F. Lim, Q. Li, Y. S. Ong, S. A. Khan, and K. Hippalgaonkar, Evolution-guided Bayesian Optimization for Constrained Multi-objective Optimization in Self-driving Labs, *npj Computational Materials* **10**, 104 (2024).
  - [101] B. Ranković and P. Schwaller, GOLLuM: Gaussian Process Optimized LLMs – Reframing LLM Finetuning through Bayesian Optimization, *arXiv* [10.48550/arXiv.2504.06265](https://arxiv.org/abs/10.48550/arXiv.2504.06265) (2025).
  - [102] M. Adachi, B. Planden, D. A. Howey, M. A. Osborne, S. Orbell, N. Ares, K. Muandet, and S. L. Chau, Looping in the Human: Collaborative and Explainable Bayesian Optimization, *Advances in Neural Information Processing Systems* [10.48550/arXiv.2310.17273](https://arxiv.org/abs/10.48550/arXiv.2310.17273) (2024).
  - [103] J. R. Deneault *et al.*, Preferential Bayesian Optimization Improves the Efficiency of Printing Objects with Subjective Qualities, *Digital Discovery* [10.1039/D4DD000320A](https://doi.org/10.1039/D4DD000320A) (2025).
  - [104] A. Ishii, R. Kamijyo, A. Yamanaka, and A. Yamamoto, BOXVIA: Bayesian optimization executable and visualizable application, *SoftwareX* **18**, 101019 (2022).
  - [105] R. Tamura, K. Tsuda, and S. Matsuda, NIMS-OS: an automation software to implement a closed loop between artificial intelligence and robotic experiments in materials science, *Science and Technology of Advanced Materials: Methods* **3**, 2232297 (2023).
  - [106] R. Arróyave, D. Khatamsaz, B. Vela, R. Couperthwaite, A. Molkeri, P. Singh, D. D. Johnson, X. Qian, A. Srivastava, and D. Allaire, A Perspective on Bayesian Methods Applied to Materials Discovery and Design, *MRS Communications* **12**, 1037 (2022).